

Polyhedral Network Aware Task Scheduling

Martin Kong

Brookhaven National Laboratory

October 14th, 2017
9th Annual Concurrent Collections Workshop
College Station, TX, USA



BROOKHAVEN
NATIONAL LABORATORY



U.S. DEPARTMENT OF
ENERGY

Introduction

- ▶ Context: Shared and distributed memory computers, or any hardware where the underlying network can be modeled
- ▶ We introduce different virtual network topologies that allow to model concepts such as direction and distance in order to produce a task schedule
- ▶ As an application, we choose the task isolation problem: how to schedule tasks which access shared resources
- ▶ Networks and hierarchies are pervasive in computer science, i.e. memory systems, computation graphs, programs
- ▶ Concrete instances of this problem are:
 - ▶ False sharing
 - ▶ Placing tasks close or nearby to fixed shared resources
 - ▶ Scheduling access of R/W resources (e.g. a file or freshly created data block)
 - ▶ Job scheduling and accessing a file system (explicit hierarchy) with fast and slow storage systems

Motivation

- ▶ Large scale applications require substantial domain, algorithmic and hardware/software knowledge
- ▶ Current and emerging technologies pose an enormous challenge in terms of performance portability and user productivity
- ▶ Data movement and long latencies are strong limiting factors for performance
- ▶ Knowledge of the underlying network topology, even in shared-memory machines, is essential to performance tuning
- ▶ Diverse background of users calls for abstractions that allow to easily switch between prescriptive and descriptive programming models and methods
- ▶ New abstractions are necessary in order to exploit and integrate network-aware compiler optimizations which:
 - Avoid communication
 - Perform efficient communication: minimize synchronization and data movement
 - Schedule tasks around data
 - Transfer domain knowledge to the compiler

A Look into the (near) Future

Feature	Titan	Summit
Application Performance	Baseline	5-10x Titan
Number of Nodes	18,688	~4,600
Node Performance	1.4 TF	> 40 TF
Feature	32 GB DDR3 + 6 GB GDDR5	512 GB DDR4 + HBM
NV Memory per Node	0	1600 GB
Total System Memory	710 TB	10 PB DDR4 + HBM + Non-Volatile
System Interconnect (Injection Bandwidth)	Gemini (6.4 GB/s)	Dual Rail EDR-IB (23 GB/s)
Interconnect Topology	3D Torus	Non-blocking Fat Tree
Processor	1 AMD Opteron + 1 NVIDIA Kepler	2 IBM Power9 + 6 NVIDIA Volta
File System	32 PB, 1 TB/s, Lustre	250 PB, 2.5 TB/s GPFS
Peak Power	9 MW	15 MW

Context

- ▶ Previously considered data dependences and “performance dependences” (e.g. additional dependences that affect scheduling, directly or indirectly)
- ▶ We leverage polyhedral tools to model different network topologies
- ▶ Goal: composability of networks
- ▶ Polyhedral compilation frameworks leverage lexicographic minimization
- ▶ Optimization of Manhattan-distance type problems are not immediately modelable
- ▶ In this work, we provide abstractions for modeling different types of network topologies
- ▶ We follow an iterative optimization process

Contributions

CnC¹

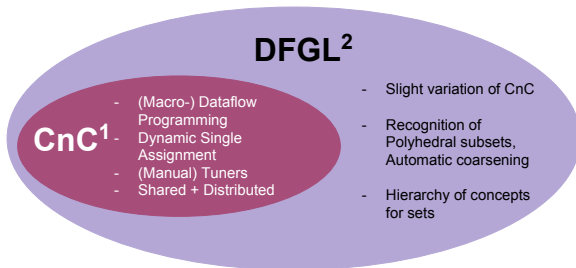
- (Macro-) Dataflow Programming
- Dynamic Single Assignment
- (Manual) Tuners
- Shared + Distributed

¹ Budimlic et. al, “Concurrent Collections”, Scientific Programming, 2010

² A. Sbirlea, L.-N Pouchet and V. Sarkar, “DFGR: an intermediate graph representation for macro-dataflow programs”, DFM, 2014

Kong, Pouchet, Sadayappan, Sarkar, “PIPES: A Language and Compiler for Task-Based Programming on Distributed-Memory Clusters”, SC, 2016

Contributions

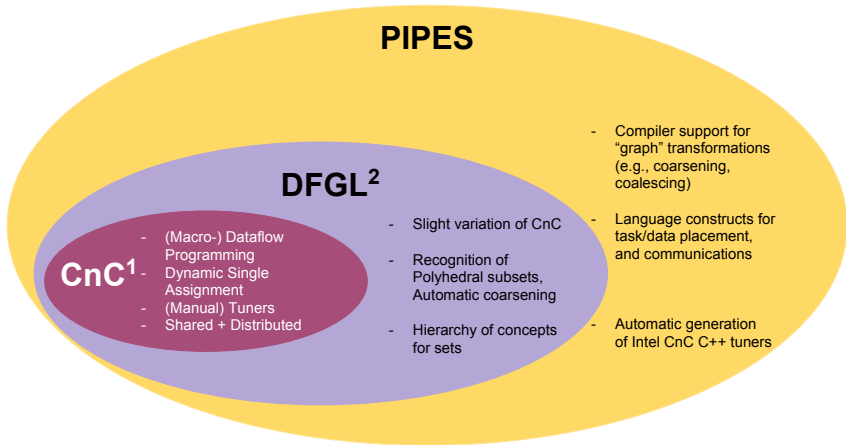


¹ Budimlic et. al, "Concurrent Collections", Scientific Programming, 2010

² A. Sbirlea, L.-N Pouchet and V. Sarkar, "DFGR: an intermediate graph representation for macro-dataflow programs", DFM, 2014

Kong, Pouchet, Sadayappan, Sarkar, "PIPES: A Language and Compiler for Task-Based Programming on Distributed-Memory Clusters", SC, 2016

Contributions

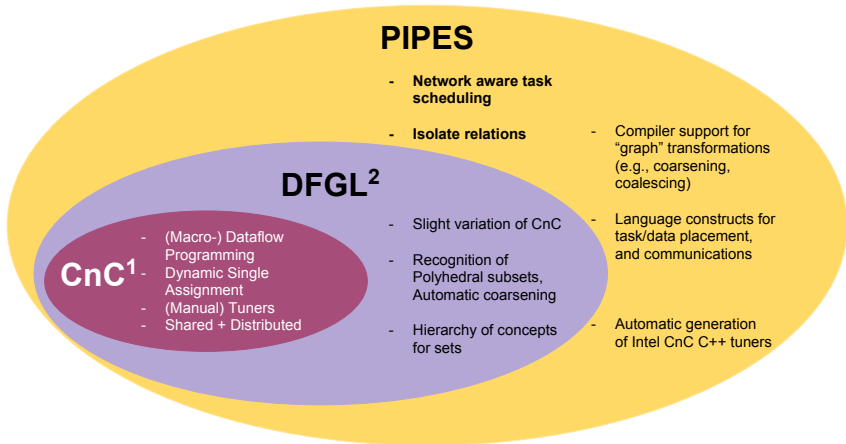


¹ Budimlic et. al, "Concurrent Collections", Scientific Programming, 2010

² A. Sbirlea, L.-N Pouchet and V. Sarkar, "DFGR: an intermediate graph representation for macro-dataflow programs", DFM, 2014

Kong, Pouchet, Sadayappan, Sarkar, "PIPES: A Language and Compiler for Task-Based Programming on Distributed-Memory Clusters", SC, 2016

Contributions



¹ Budimlic et. al, "Concurrent Collections", Scientific Programming, 2010

² A. Sbirlea, L.-N Pouchet and V. Sarkar, "DFGR: an intermediate graph representation for macro-dataflow programs", DFM, 2014

Kong, Pouchet, Sadayappan, Sarkar, "PIPES: A Language and Compiler for Task-Based Programming on Distributed-Memory Clusters", SC, 2016

PIPES Language Features

- ▶ Language features are task-centric
- ▶ Virtual topologies
- ▶ Task placement
- ▶ Data placement
- ▶ Data communication (pull or push communication model)

MatMul in PIPES

```
1 Parameter N, P;
2 // Define data collections
3 [float* A:1..N,1..N];
4 [float* B:1..N,1..N];
5 [float* C:1..N,1..N,1..N+1];
6 // Task prescriptions
7 env :: (MM:1..N,1..N,1..N);
8 // Input/Output:
9 env -> [A:1..N,1..N];
10 env -> [B:1..N,1..N];
11 env -> [C:1..N,1..N,1];
12 [C:1..N,1..N,N+1] -> env;
13 // Task dataflow
14 [A:i,k], [B:k,j], [C:i,j,k] -> (MM:i,j,k) -> [C:i,j,k+1];
```

Figure: PIPES Matrix Multiplication

Language Construct Summary

Name	Syntax
Region	<code>regname = { tuple : constraints } properties</code>
Prescription Relation	<code>env :: (task : regname)</code>
Data-Flow Relations	<code>[input_instances] -> (task_instance) -> [output_instance]</code>
Virtual Topology	<code>Topology topo_name = { sizes=[parameter_list] }</code>
Affinity Mapping	<code>(task : tuple) @ [[topo : tuple]]</code>
Communication	<code>[item: tuple] @ (task1 : tuple) => (task2 : tuple)</code>
Scheduling	<code>(task1 : tuple) -> (task2 : tuple)</code> <code>(task1 : tuple) ~> (task2 : tuple)</code>

Table: PIPES Language Constructs

Virtual Topologies and Task Mapping

- ▶ Virtual topologies (VTs) represent the logical underlying computer grid/cluster
- ▶ Each element in the set is a processor
- ▶ Requires a logical-to-physical mapping

```
1 // 2D topology, no more than 256  
  x256 processors  
2 Parameter P : 1..256;  
3 Topology Topo2D = {  
4   sizes=[P,P];  
5   cores=[i,j] : { 0 <= i < P, 0  
                     <= j < P};  
6 };
```

Virtual Topologies and Task Mapping

- ▶ Mappings of tasks to elements in the topology
- ▶ Task (instance) will execute on the processor it is mapped to
- ▶ Always enforced by run-time
- ▶ Requires the topology to be defined
- ▶ Maps directly to the **compute_on** tuner

```
1 (task:tag-set) @ Topo2d(point);
```

Overall Approach






- ▶ Network topology abstractions
- ▶ Map task instances to virtual topology via Affinity Tasks Mappings (ATM)
- ▶ Introduce **isolate relations** and **isolate dependences**
- ▶ Compute in closed for the possible task interleavings
- ▶ Filter / prune unwanted tasks orderings from the closed form
- ▶ For each point in the closed form, iteratively compute its cost
- ▶ Apply lexicographic minimization to determine the best execution order

Network Topologies

- ▶ PIPES uses an abstraction that allow to pin task instances to processing elements / cores; main purpose is to determine if tasks execute on the same location; then generate Intel CnC++ tuners
- ▶ In this work: Implemented two different topologies: mesh and fat tree
- ▶ Pending to implement: torus / meshes + wraparound, hypercubes, hierarchies
- ▶ Each network type requires different set of abstractions
- ▶ Abstractions allow modeling concepts such as dimension, distance or direction
- ▶ In the future, want to pursue composing two or more topology types, so as to faithfully represent upcoming HPC and Data Analytics clusters

Driving Example

- ▶ Consider a 4x4 mesh and 4 tasks (A,B,C,D)
- ▶ What if A has to start the computation?
- ▶ What if any task can start the computation?
- ▶ What if wrap-around communication is allowed?
- ▶ In general, avoiding ping-pong-ing around



Network Topologies: N-dimensional Meshes

Key: idea: translate task tuple to network coordinates, then compute distance between network coordinates.

Example:

- 1 Compute task topology coordinates:

$A=G[0,0]; \quad B=G[3,3];$

$C=G[0,3]; \quad D=G[1,0]$

- 2 Compute task-to-task topology directions:

$\text{Dir}(A,B) = [3,3]; \quad \text{Dir}(B,C) = [0,-3];$

$\text{Dir}(C,D) = [-3,1]$

- 3 Compute "positive directions":

$+\text{Dir}(\text{Dir}(A,B)) = [3,3];$

$+\text{Dir}(\text{Dir}(B,C)) = [0,3];$


$+\text{Dir}(\text{Dir}(C,D)) = [3,1];$

- 4 Compute task-to-task topology distance:

$\text{Dist}(A,B) = \text{MultiplexAddMap}([3,3]) = 6$

$\text{Dist}(B,C) = \text{MultiplexAddMap}([0,3]) = 3$


$\text{Dist}(C,D) = \text{MultiplexAddMap}([3,1]) = 4$



Network Topologies: Fat Trees

- ▶ Distance metrics in fat tree type of network are non-linear
- ▶ Coordinate space is one dimensional
- ▶ Task tuples must be “flattened” to represent processors in the grid
- ▶ Restricted to fixed (non-parametric) network sizes
- ▶ Approach: explicitly construct a distance map from a pair of processors to a fixed non-parameteric value, i.e.

$$[[P_i] - > [P_j]] - > [distance]$$
- ▶ Does not require the **coordinate to direction** map nor the **positive direction map**
- ▶ Observe the recursive nature



	P0	P1	P2	P3	P4	P5	P6	P7
P0	0	1	2	2	4	4	4	4
P1	1	0	2	2	4	4	4	4
P2	2	2	0	1	4	4	4	4
P3	2	2	1	0	4	4	4	4
P4	4	4	4	4	0	1	2	2
P5	4	4	4	4	1	0	2	2
P6	4	4	4	4	2	2	0	1
P7	4	4	4	4	2	2	1	0

Network Topology Abstractions

- ▶ Affinity maps (task tuple to network coordinate maps)
- ▶ Coordinate to direction maps (each dimension can be +/-)
- ▶ Direction unification map: consider different routes along each dimension and direction (several disjunctions), e.g. in meshes with wraparound
- ▶ Direction to distance maps: make directions positive
- ▶ Cumulative distance maps: "Multiplex Add Map"

Building the Closed Form of the Exploration Space

- ▶ Compact and closed form for encoding task orderings at the coarse level
- ▶ Assign to each task a fixed integer
- ▶ Start with the T^T potential task orderings (T : number of **lexical** tasks), fixed and known at compile time
- ▶ The exploration space consist of the set of points in a T-dimensional tuple space.
- ▶ Add bounding constraints: $\vec{t} = (t_1, t_2, \dots, t_T), \forall t_i \in [1..T]$
- ▶ Add constraints to remove meaningless points, e.g. (1,1,1) which represents the same task being executed: $\sum t_i = T \times (T + 1)/2$
- ▶ Add constraints to enforce data dependences, i.e. if task 1 is a dependence of task 2 then: $t_i = 1$ and $t_j = 2$, for some $j > i$
- ▶ Each point in the set represents a full execution path (e.g. $3 \rightarrow 1 \rightarrow 4 \rightarrow 2$)

Pruning the Scheduling Space

We consider three types of pruning constraints that are added to the complete space, e.g. $(t_1, t_2, \dots, t_T), t_i \in [1..T]$

- ▶ Data dependence edges: “A(1) must to execute before B(2)”
- ▶ Scheduling edges (for performance): “We want A(1) to execute before B(2)”
- ▶ **Isolation edges**: “We want either A(1) before B(2) or B(2) before A(1)”

Several disjunctions might be required to model the constraints, in particular for the **isolation edges**

A few more considerations:

- ▶ In theory, we could still have meaningless points, e.g. $(1, 1, 4, 4) = 10$
- ▶ In practice, these would be mostly naturally pruned by the above constraints
- ▶ To be safe, we do a manual check before the next stage to guarantee that all entries in a vector are distinct

Iteratively Computing Path Costs

```

1: space  $\leftarrow$  compute_closed_form (T,program)
2: for each  $\vec{t} = (t_1, t_2, \dots, t_N) \in \text{space}$  do
3:   path( $\vec{t}$ )  $\leftarrow$  0
4:   for each  $(t_i, t_{i+1}) \in (t_1, t_2, \dots, t_N)$  do
5:     Fetch task domains  $T(t_i)$  and  $T(t_{i+1})$ 
6:     Fetch task affinity maps:  $AM(t_i)$  and  $AM(t_{i+1})$ 
7:     Build synchronization map: sync_map  $\leftarrow T(t_i) \rightarrow T(t_{i+1})$ 
8:     Intersect dependences  $T(t_i) \Rightarrow T(t_{i+1})$  with sync_map
9:     Compute processor synchronization map (PSM): PSM
        $\leftarrow AM(t_i)^{-1} \circ \text{sync} \circ AM(t_{i+1})$ 
10:    Compute “processor coordinate difference” (PCD) with ISL
       deltas_map (PSM)
11:    Apply network specific coordinate-to-distance map to PCD
12:    edge(i)  $\leftarrow$  quasi_polynomial_sum(PSM)
13:    path( $\vec{t}$ )  $\leftarrow$  path( $\vec{t}$ ) + edge(i)
14:  end for
15:  M  $\leftarrow$  M  $\cup$  path( $\vec{t}$ )
16: end for
17: result  $\leftarrow$  lexmin(M)

```

Applications of this Technique

- ▶ Minimize synchronization latency among tasks
- ▶ Pre-optimization pass that affects the overall program prior to applying high-level loop transformations that optimize for locality, akin to code motion in for loops
- ▶ Allows to establish order among tasks that should be executed in a non-concurrent fashion:
 - 1 User provides **isolate relations**, e.g. “ $A \sim || B$ ”, another class of dependence that will be enforce semantic orderings, both at the compiler and runtime level
 - 2 Compiler “decides the direction” based on network distance/latency
 - 3 Isolate relations are then promoted to “isolate dependences” and fed to some other scheduler

Applications of this Technique

- ▶ This technique has the potential to reduce the runtime scheduling overhead by substantially narrowing down the scheduling options
- ▶ Autotuning: coupled with auto-generated task mappings, allows to determine suitable mappings and task schedules
- ▶ Applicable to several task parallel and data-flow runtimes (CnC!!)

Restrictions of the Approach

- ▶ Very computational expensive, even in non-parametric cases
- ▶ Limited to ~ 10 tasks (millions of possible interleavings) and taking ~ 10 min
- ▶ Disallow edges that induce loops in the graph
- ▶ Some networks cannot be modeled with affine parametrics constraints, resort to use large fixed values and bound network parameters by the context

Experimental Setup

- ▶ OS: Mac Sierra
- ▶ 3.5 GHz Intel Core i7
- ▶ Memory: 16 GB 2133 MHz
- ▶ Compiler: Clang++ Apple LLVM version 8.1.0 (clang-802.0.42)
- ▶ Barvinok 40
- ▶ ISL 18.0

Will show some preliminary compilation results for 2-D meshes (with fixed and parametric task domains) and fat-trees.

Mesh Results: Fixed

Test	Tasks	Deps	Factorial	Legal	Semi-legal	Complete	Time (sec)
1	2	1	2	1	2	4	0.021
2	2	0	2	2	2	4	0.059
3	2	0	2	2	2	4	0.101
4	3	2	6	2	4	27	0.153
5	5	4	120	2	130	3125	0.205
6	2	1	2	2	2	4	0.027
7	3	2	6	2	4	27	0.161
8	3	3	6	4	4	27	0.069
9	3	3	6	4	4	27	0.094

Space exploration size (T : number of lexical tasks):

- ▶ Factorial: $T!$
- ▶ Legal: Final exploration space
- ▶ Semi-Legal: Exploration space before adding dependence edges
- ▶ Complete: T^T exploration space

Mesh Results: Parametric

Test	Tasks	Deps	Factorial	Legal	Semi-legal	Complete	Time (sec)
01	2	0	2	2	2	4	0.056
02	3	2	6	2	7	27	0.037
03	3	1	6	3	7	27	
04	4	1	24	17	44	256	
05	5	1	120	120	381	3,125	0.1
06	5	4	120	4	381	3,125	
07	6	5	720	12	4,332	46,656	
08	7	5	5,040	84	60,691	823,543	0.15
09	9	7	362,880	648	19,610,233	387,420,489	0.96
10	9	8	362,880	1	19,610,233	387,420,489	233
11	10	9	3.63E+06	2	432457640	10,000,000,000	28.8
							317.2

Space exploration size (T : number of lexical tasks):

- ▶ Factorial: $T!$
- ▶ Legal: Final exploration space
- ▶ Semi-Legal: Exploration space before adding dependence edges
- ▶ Complete: T^T exploration space

Fat Tree Results

Test	Fat Tree Max Size	Tasks	Deps	Factorial	Legal	Semi-legal	Complete	Time (min + sec)
1	16	2	1	2	2	2	4	0m0.077s
2	32	2	1	2	2	2	4	0m0.167s
3	64	2	1	2	2	2	4	0m0.495s
4	128	2	1	2	2	2	4	0m1.583s
5	256	2	1	2	2	2	4	0m5.556s
6	512	2	1	2	2	2	4	0m21.163s
7	1024	2	1	2	2	2	4	1m21.719s
8	2048	2	1	2	2	2	4	5m31.795s
9	256	5	3	120	12	130	3125	2m5.986s
10	256	5	4	120	34	130	3125	5m58.886s
11	256	5	6	120	130	130	3125	22m57.720s

Space exploration size (T : number of lexical tasks):

- ▶ Factorial: $T!$
- ▶ Legal: Final exploration space
- ▶ Semi-Legal: Exploration space before adding dependence edges
- ▶ Complete: T^T exploration space

Future Work

- ▶ Complete implementation of torus networks, hypercubes and explicit hierarchies
- ▶ Enable composability of network topology abstractions
- ▶ Integrate resource features into topology abstraction:
 - To model **super nodes** that have access near to accelerators or to memory nodes
 - Specialized resource types e.g. streaming nodes, compute intensive nodes, memory nodes, storage nodes, etc
- ▶ Potential direction: focus on file access scheduling (e.g. in HPF5) by using the data file and block structure to construct the program graph
- ▶ Complete integration to PIPES compiler
- ▶ Leverage the newly introduced constructs to efficiently implement and perform communication patterns such as:
 - All to all communication
 - Multicast and broadcast
 - Nearest neighbor type of communication
- ▶ Improve the scalability of the technique: Perform cuts on the graph, search for “bridge tasks” and “articulation tasks”

Das Ende

Thanks for listening

Questions?