

# Beyond Array-based CnC Apps

CnC Workshop, October 14th, 2017

Kath Knobe (Rice University)

Zoran Budimlić (Rice University)



# Current state

- Consider a presentation at a large conference:

```
X(row, col, iter) = Func(  
    X(row, col, iter-1),  
    X(row+1, col, iter-1),  
    X(row, col+1, iter-1),  
    X(row-1, col, iter-1),  
    X(row, col-1, iter-1) )
```

- Everyone knows what this means
  - Domain experts can read & write
  - Optimizers can optimize
  - Runtimes can execute
- The array-based community is well served by + and -

# +1 and -1 are about semantics of arrays

- They are not about how arrays are stored
  - e.g., row-major, col-major, space-filling curves, double buffering the iterations, ...
- This semantic level communication in array-based apps used:
  - In domain-specs
  - In analyzers
  - In optimizers
  - In runtime

# Tree-based apps are everywhere

- Tree-based applications might use pointers
- Runtimes can execute but ...
  - Harder to write
  - Harder to read
  - Harder to analyze
  - Harder to optimize

# We need to create a new situation for trees

The goal is

- To address the semantics of trees
  - not about how trees are stored
- So the tree-based community is also well served
  - Domain experts can write
  - Domain experts can read
  - Optimizers can analyze and optimize
  - Runtimes can execute

# Basic idea 1: Step vs dependence functions

Current thinking:

- Steps
  - CnC knows about the role of steps
  - Steps and domain-specific libraries of steps are (will be) reusable among set of domain-specific apps
- Dependence functions
  - CnC knows about the role of dependence functions
  - Dependence functions and domain-specific libraries of dependence functions will be reusable among set of domain-specific apps
- Step code and the dependence functions from application writer are passed to the analysis/optimization/execution phases

New thinking:

- Steps
  - Unchanged
- Dependence functions
  - *Dependence functions* and domain-specific libraries of dependence functions:  
are also be reusable among set of domain-specific apps
  - CnC knows their role but doesn't know the computation

# Basic idea 1

- Current thinking: Array
  - *Array*-specific analyses and optimizations of CnC graphs:  
Can be reusable among set of array-based applications and implementations
- New thinking: Domain
  - *Domain*-specific analyses and optimizations:  
should also be reusable among set of *domain*-specific applications and implementations
  - *Tree*-specific analyses and optimizations:  
Should be reusable among set of *tree*-specific applications and implementations

# Basic idea 2

## Current thinking

- CnC knows about
  - the role of steps
  - the role of dependence functions
- Step code and the dependence functions from application writer are passed to the analysis/optimization/execution phases
- Reusable libraries (current thinking)
  - Of step codes
  - CnC graphs

## Future thinking

- Reusable domain-specific libraries also includes
  - Dependence functions

# New role: Community-expert

- Add to existing roles of domain-expert and tuning-expert
- What is a community?

Community has

- A shared set of computations (reusable steps and subgraphs)
- A shared set of data structures (arrays, trees, ...)
- A shared set of dependence functions
- A shared set of analyzes & optimizations on the community structures

Community is not

- A specific CnC application
- All CnC applications
- A specific way of storing the data structures

- CnC array-based community is already within our (subconscious) thinking

Implied: shared set of dependence functions, e.g., +1, -1, north, nextIteration, ...

- Proposal:

To provide some support for tree-based community

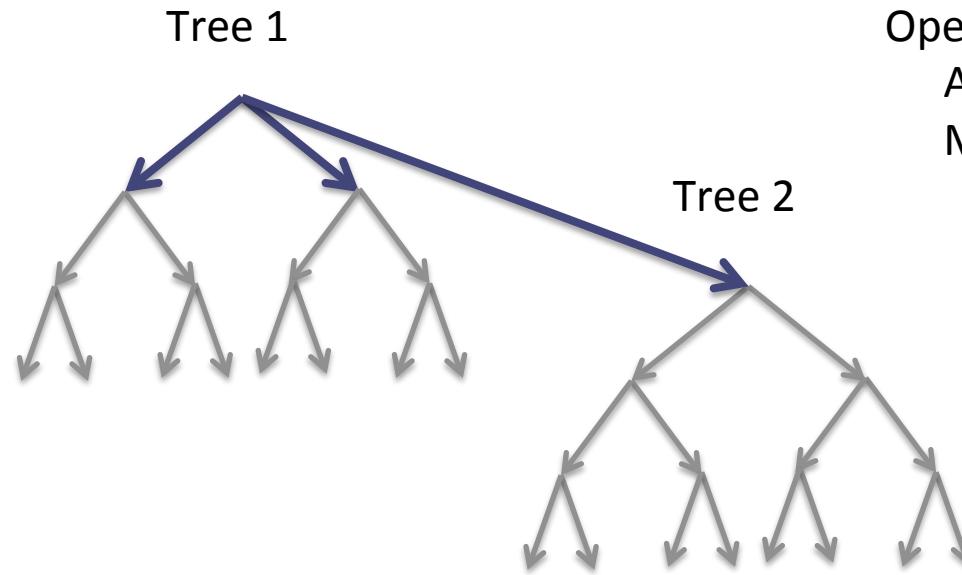
So we need to

- have a better understanding of the role of community, in general
- develop specific support for a tree-based community

# Initial work

- An approach (or a set of approaches) to tagging nodes in tree-community
  - Intuitive
  - Optimizable
- Domain-specific dependence functions applied to tree-node tags.
  - For arrays:
    - +1, -1, applied to row, col, iteration or
    - NextRow(), prevIteration()
  - For trees:
    - Parent(), firstChild(), nextSib(), nextTree(), ...
- Might want a community-specific algebra [Zoran]
  - Parent(firstChild( $n$ )) is  $n$
  - FirstChild(parent( $n$ )) may or may not be  $n$
  - NextSib(previousSib( $n$ )) = previousSib(nextSib( $n$ ))  
unless  $n$  is the first or the last child

# Sequence of operators on binary trees: 3 dependences



Operator

Assumed to be the same within a tree  
Might be different among trees

A tiling can follow

- Any one dependence
- Any pair of dependences
- All three

# (de)composition on trees

Assume:

A given tree is a given operation.

Distinct trees might have distinct operations.

	<b>Data</b>	<b>Computation</b>
Homogeneous  The children: the same code Within a tree	Arrays Trees	Like SPMD
Heterogeneous		

# Among trees: similar to heterogeneous hierarchical decomposition

Assume:

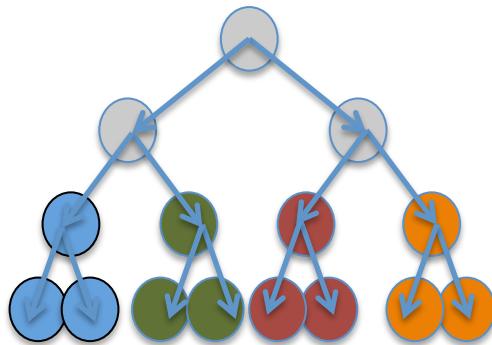
A given tree is a given operation.

Distinct trees might have distinct operations.

	Data	Computation
Homogeneous		
Heterogeneous	Like a struct  The children: different code Among trees	Like MIMD

# Example of possible tiling within a single tree.

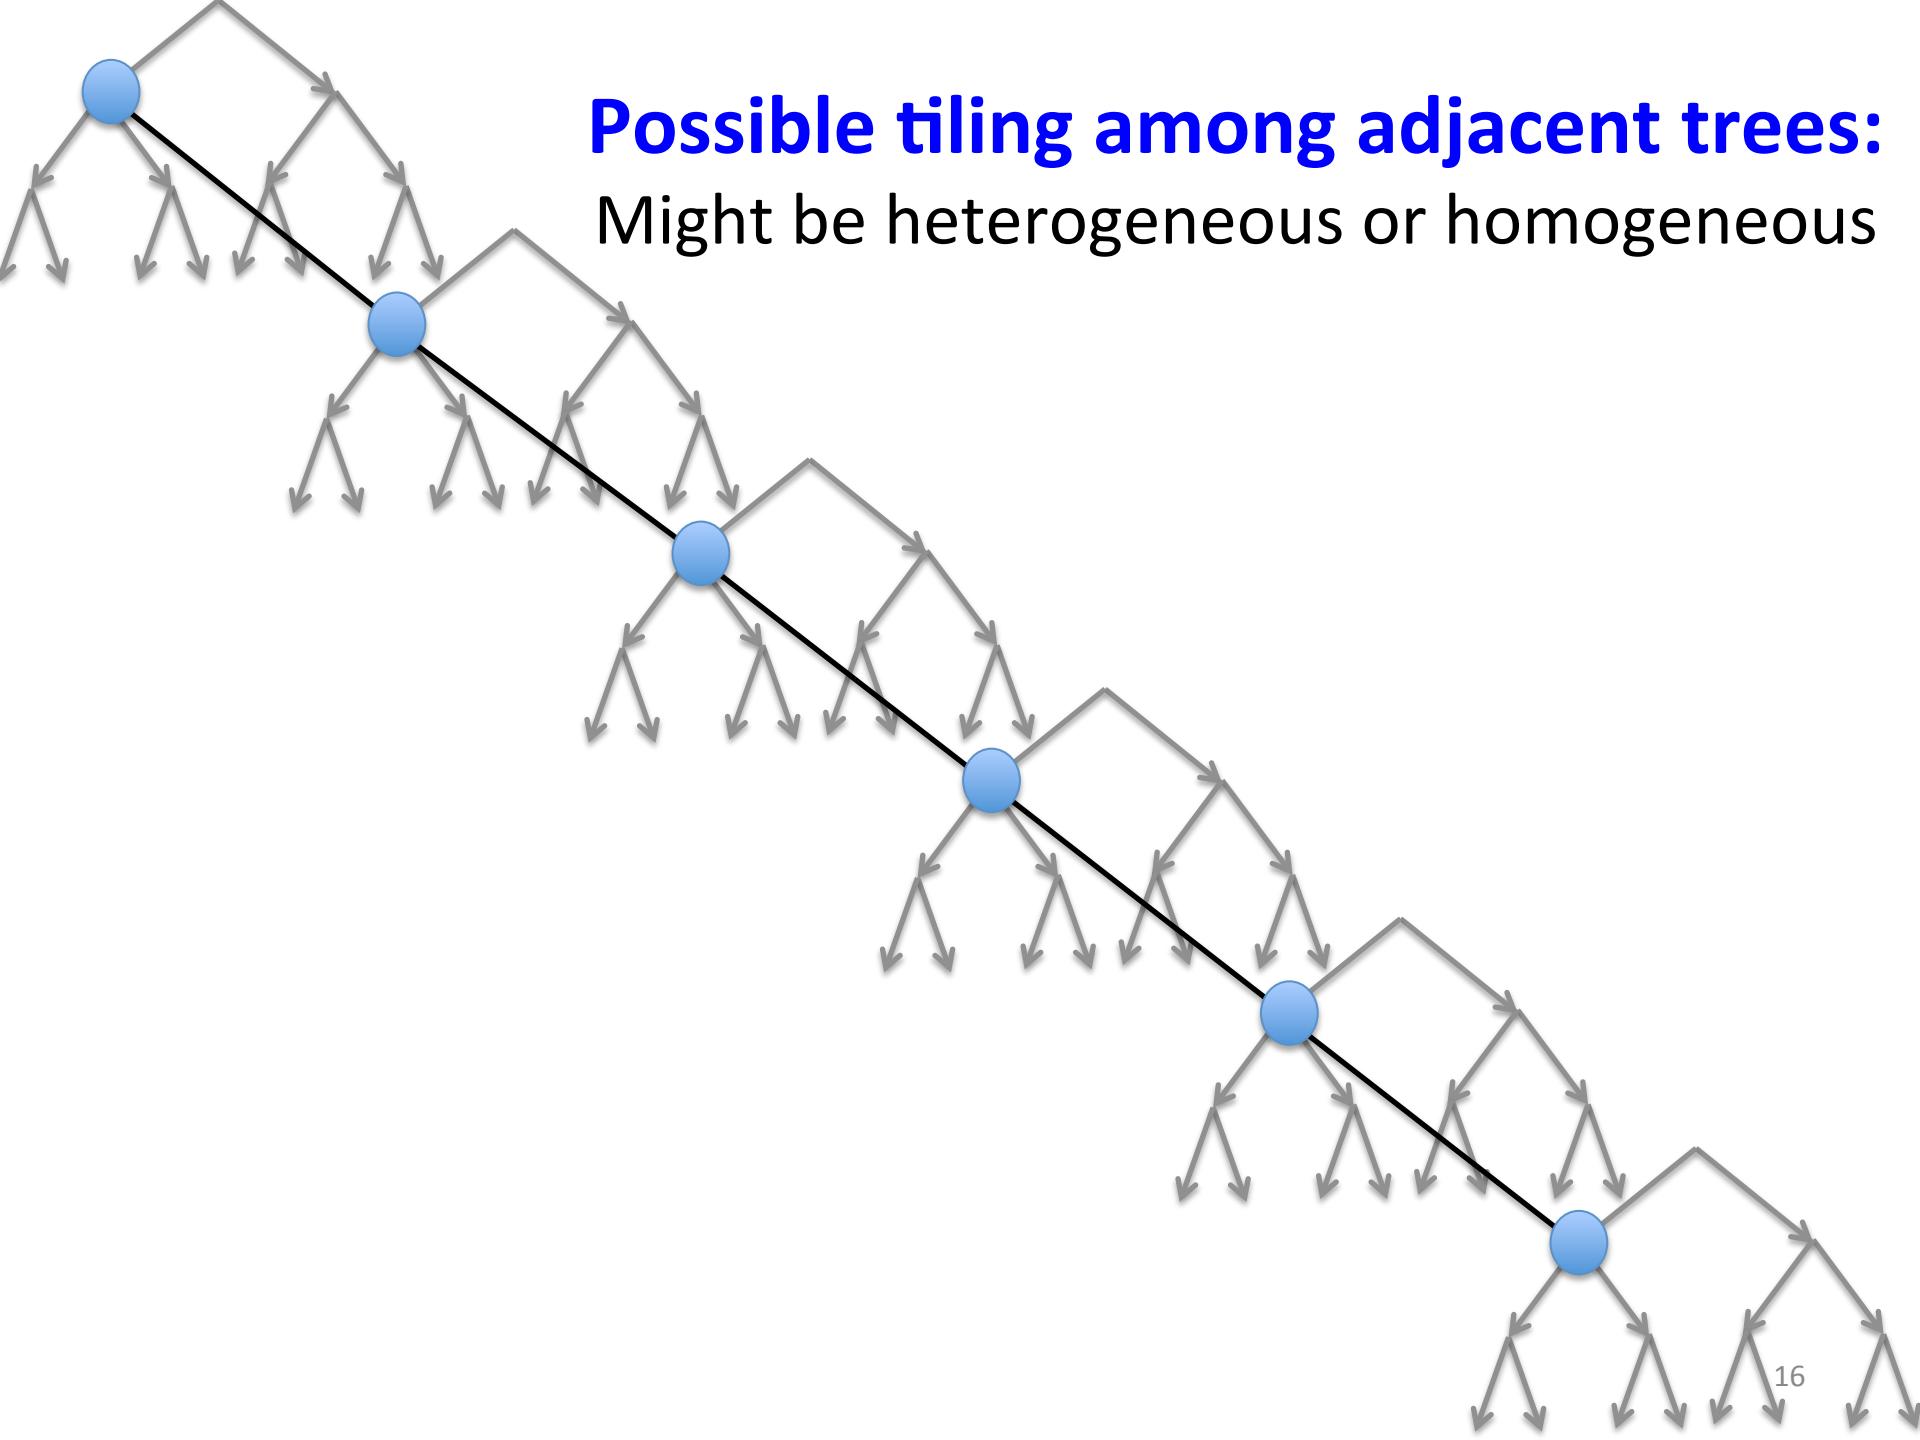
This is the tree analog of an array tile  
Within a tree => homogeneous tiles



# Heterogeneous/homogeneous

Conjecture:

- Doesn't impact tagging, e.g., `<nodesInTree, treeIDs>`
- Does impact the code for the steps
  - Homogeneous: has the code for one operation
  - Heterogeneous: has code for multiple operations



## Possible tiling among adjacent trees:

Might be heterogeneous or homogeneous

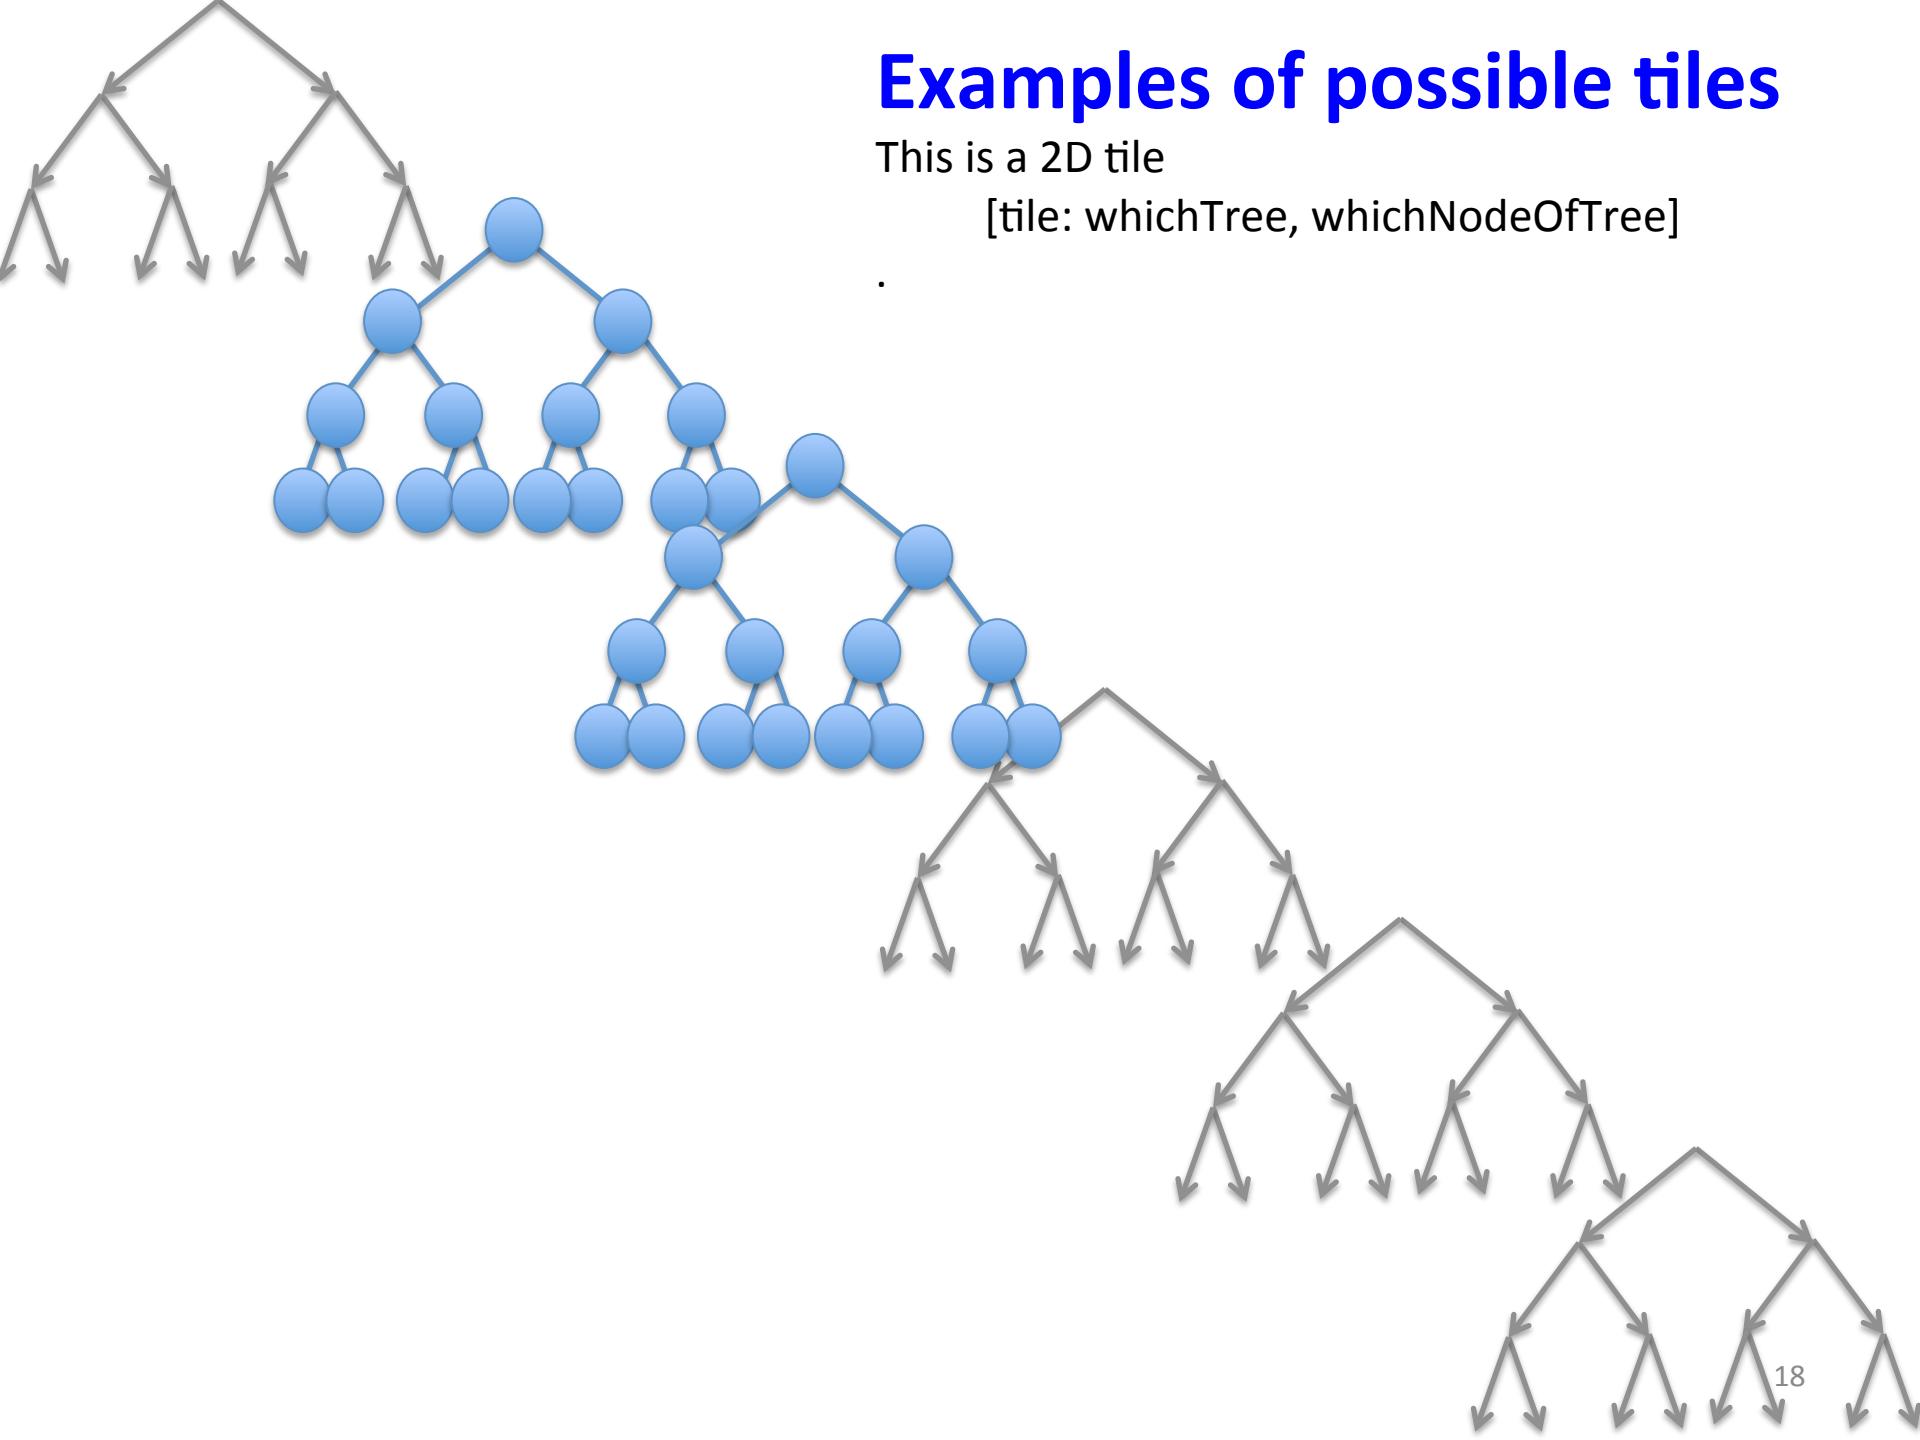
## Next

- Develop a way of tagging tree steps and items
- Agree on a set of tree dependence functions
  - E.g., LeftChild, RightChild, Parent, NextTree, PrevTree
- Develop ways of executing heterogeneous tiles
  - Alter step code
  - Static scheduling of the set of steps, each homogeneous
  - ...
- Develop a variety of walks using various tilings

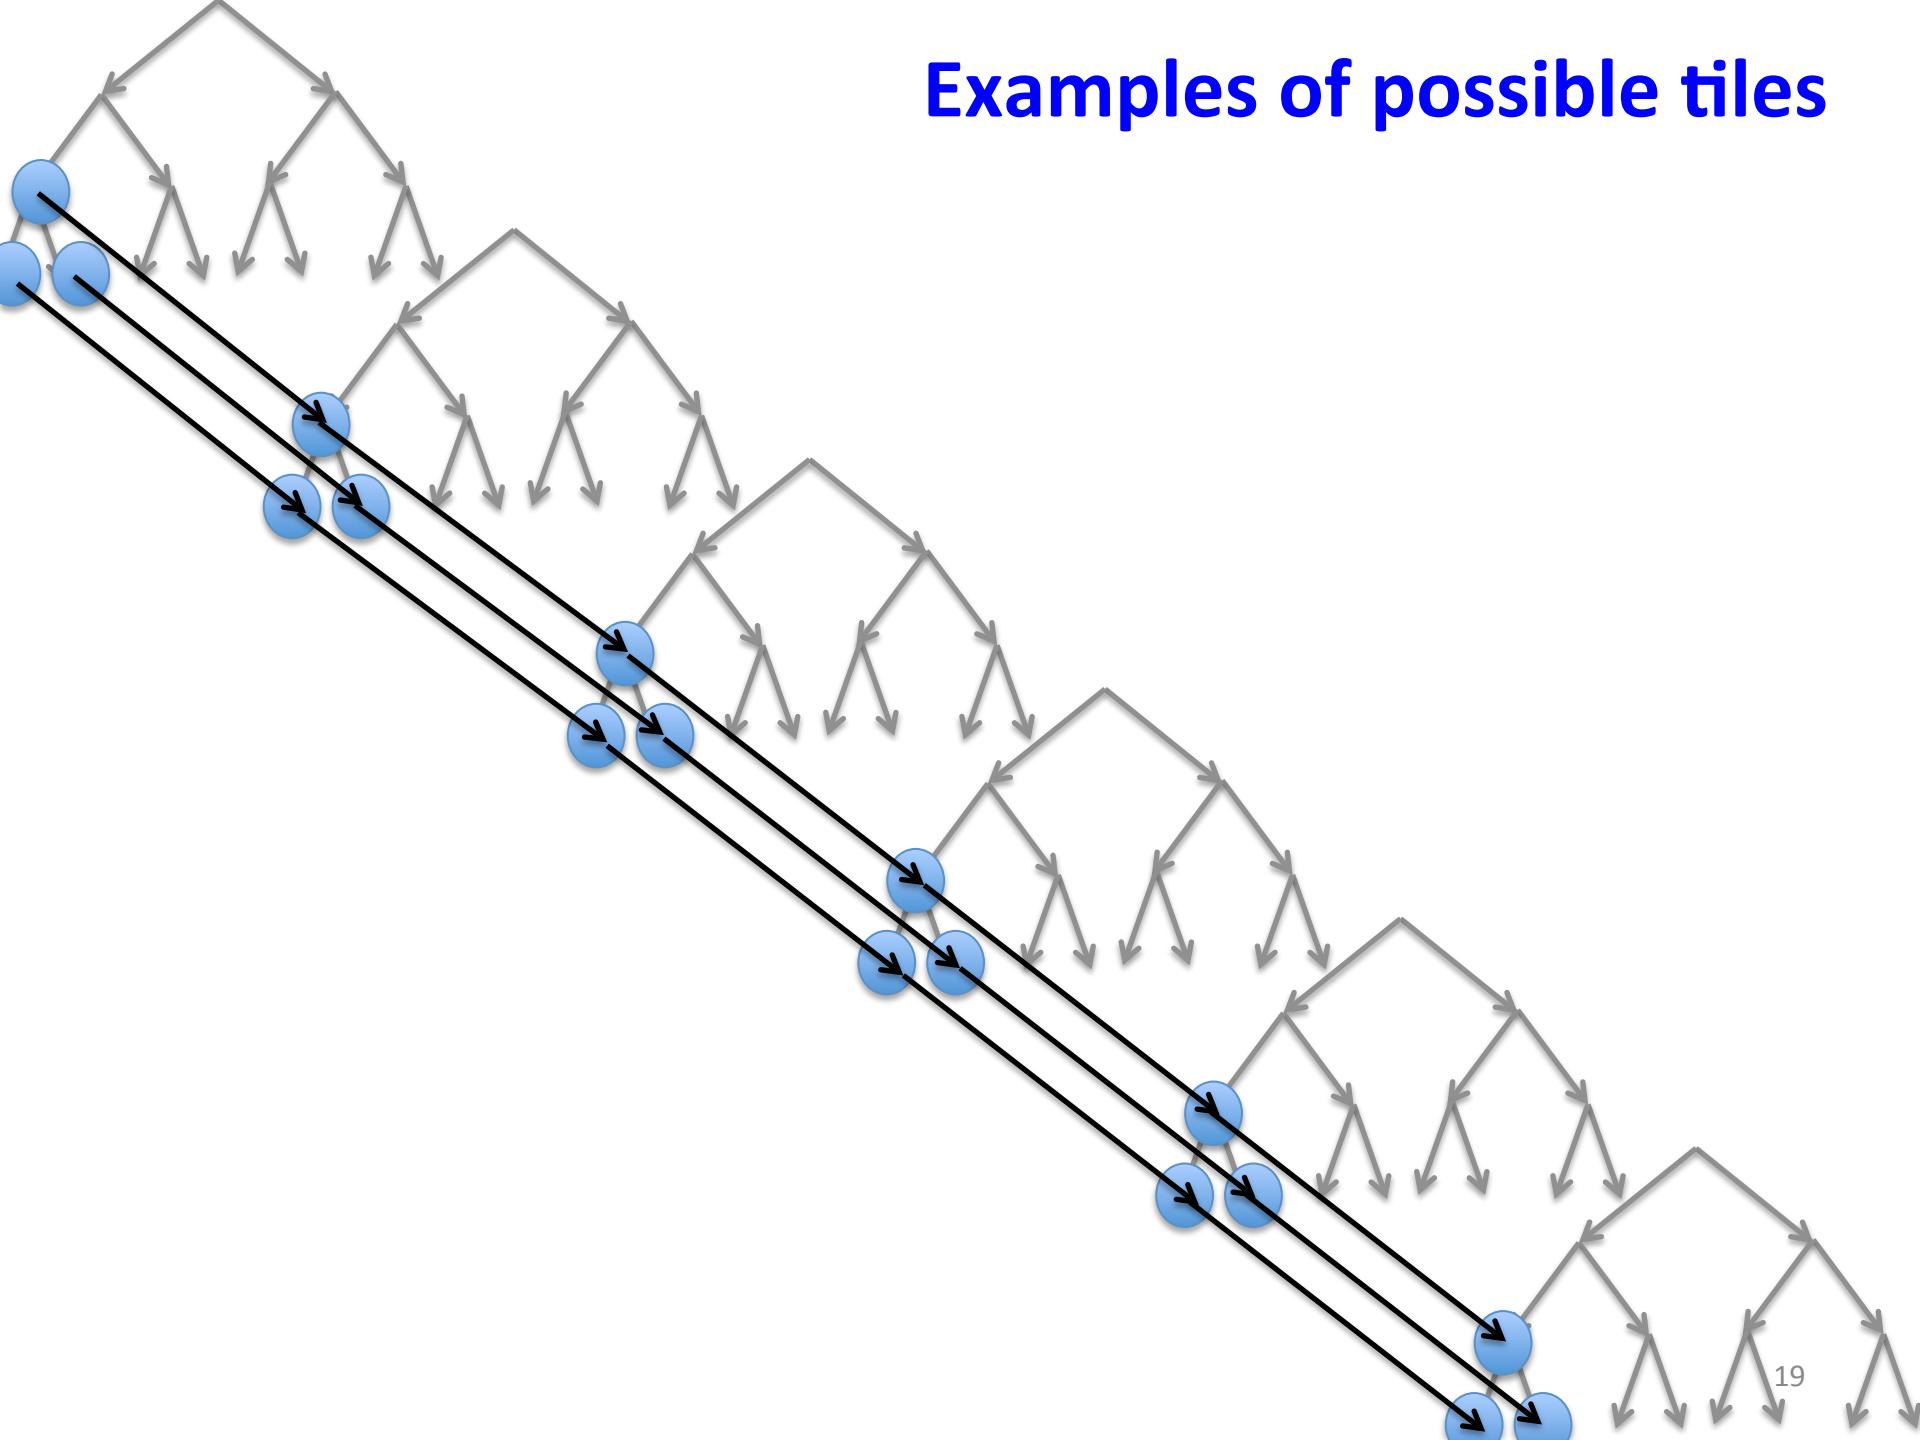
# Examples of possible tiles

This is a 2D tile

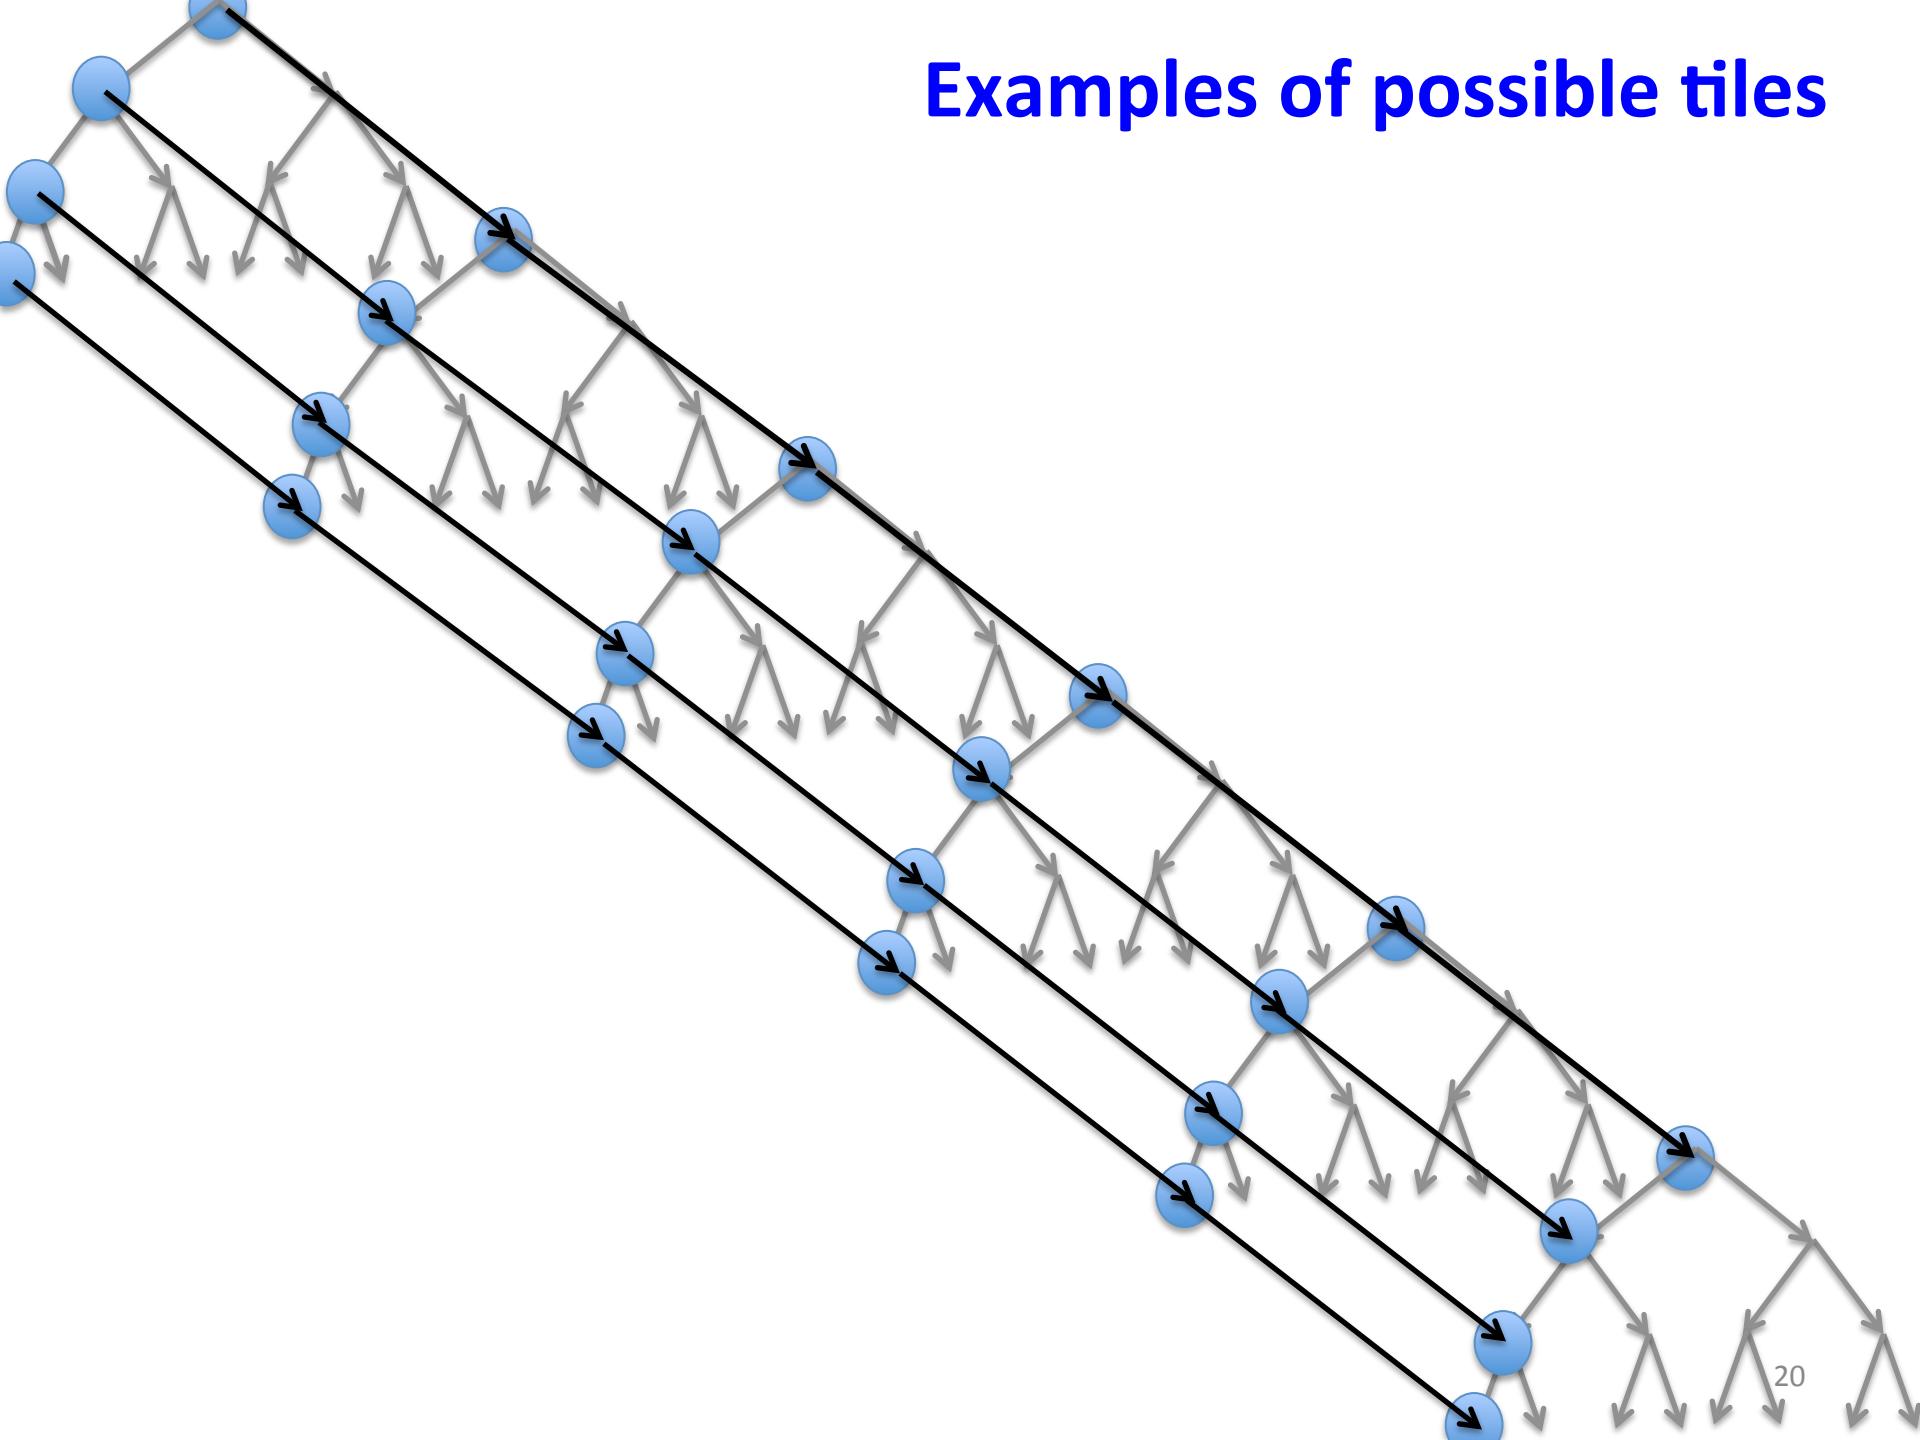
[tile: whichTree, whichNodeOfTree]



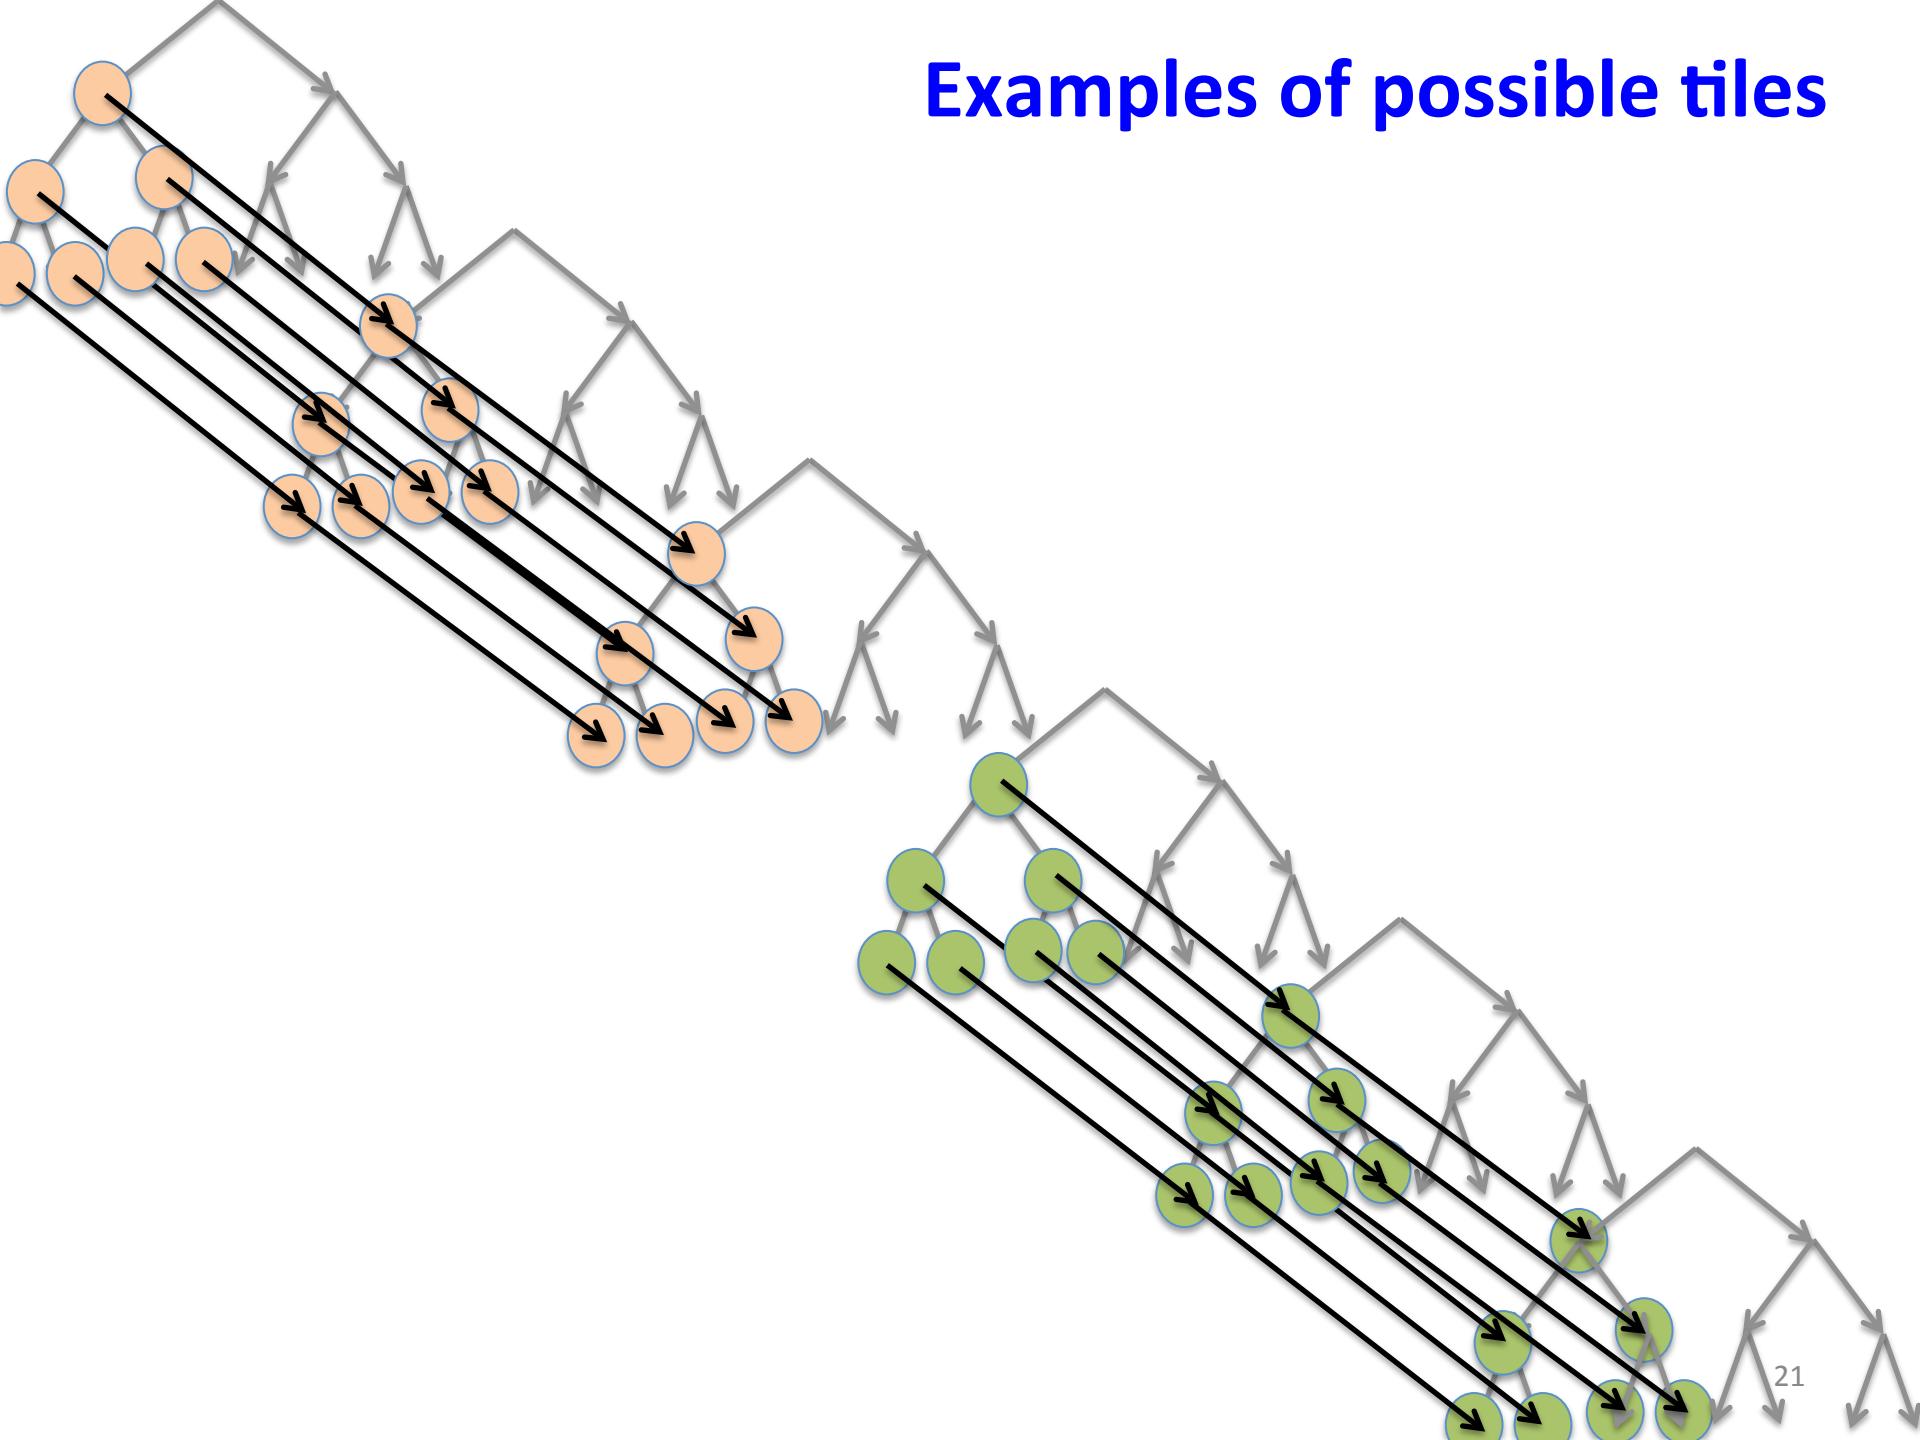
# Examples of possible tiles



# Examples of possible tiles



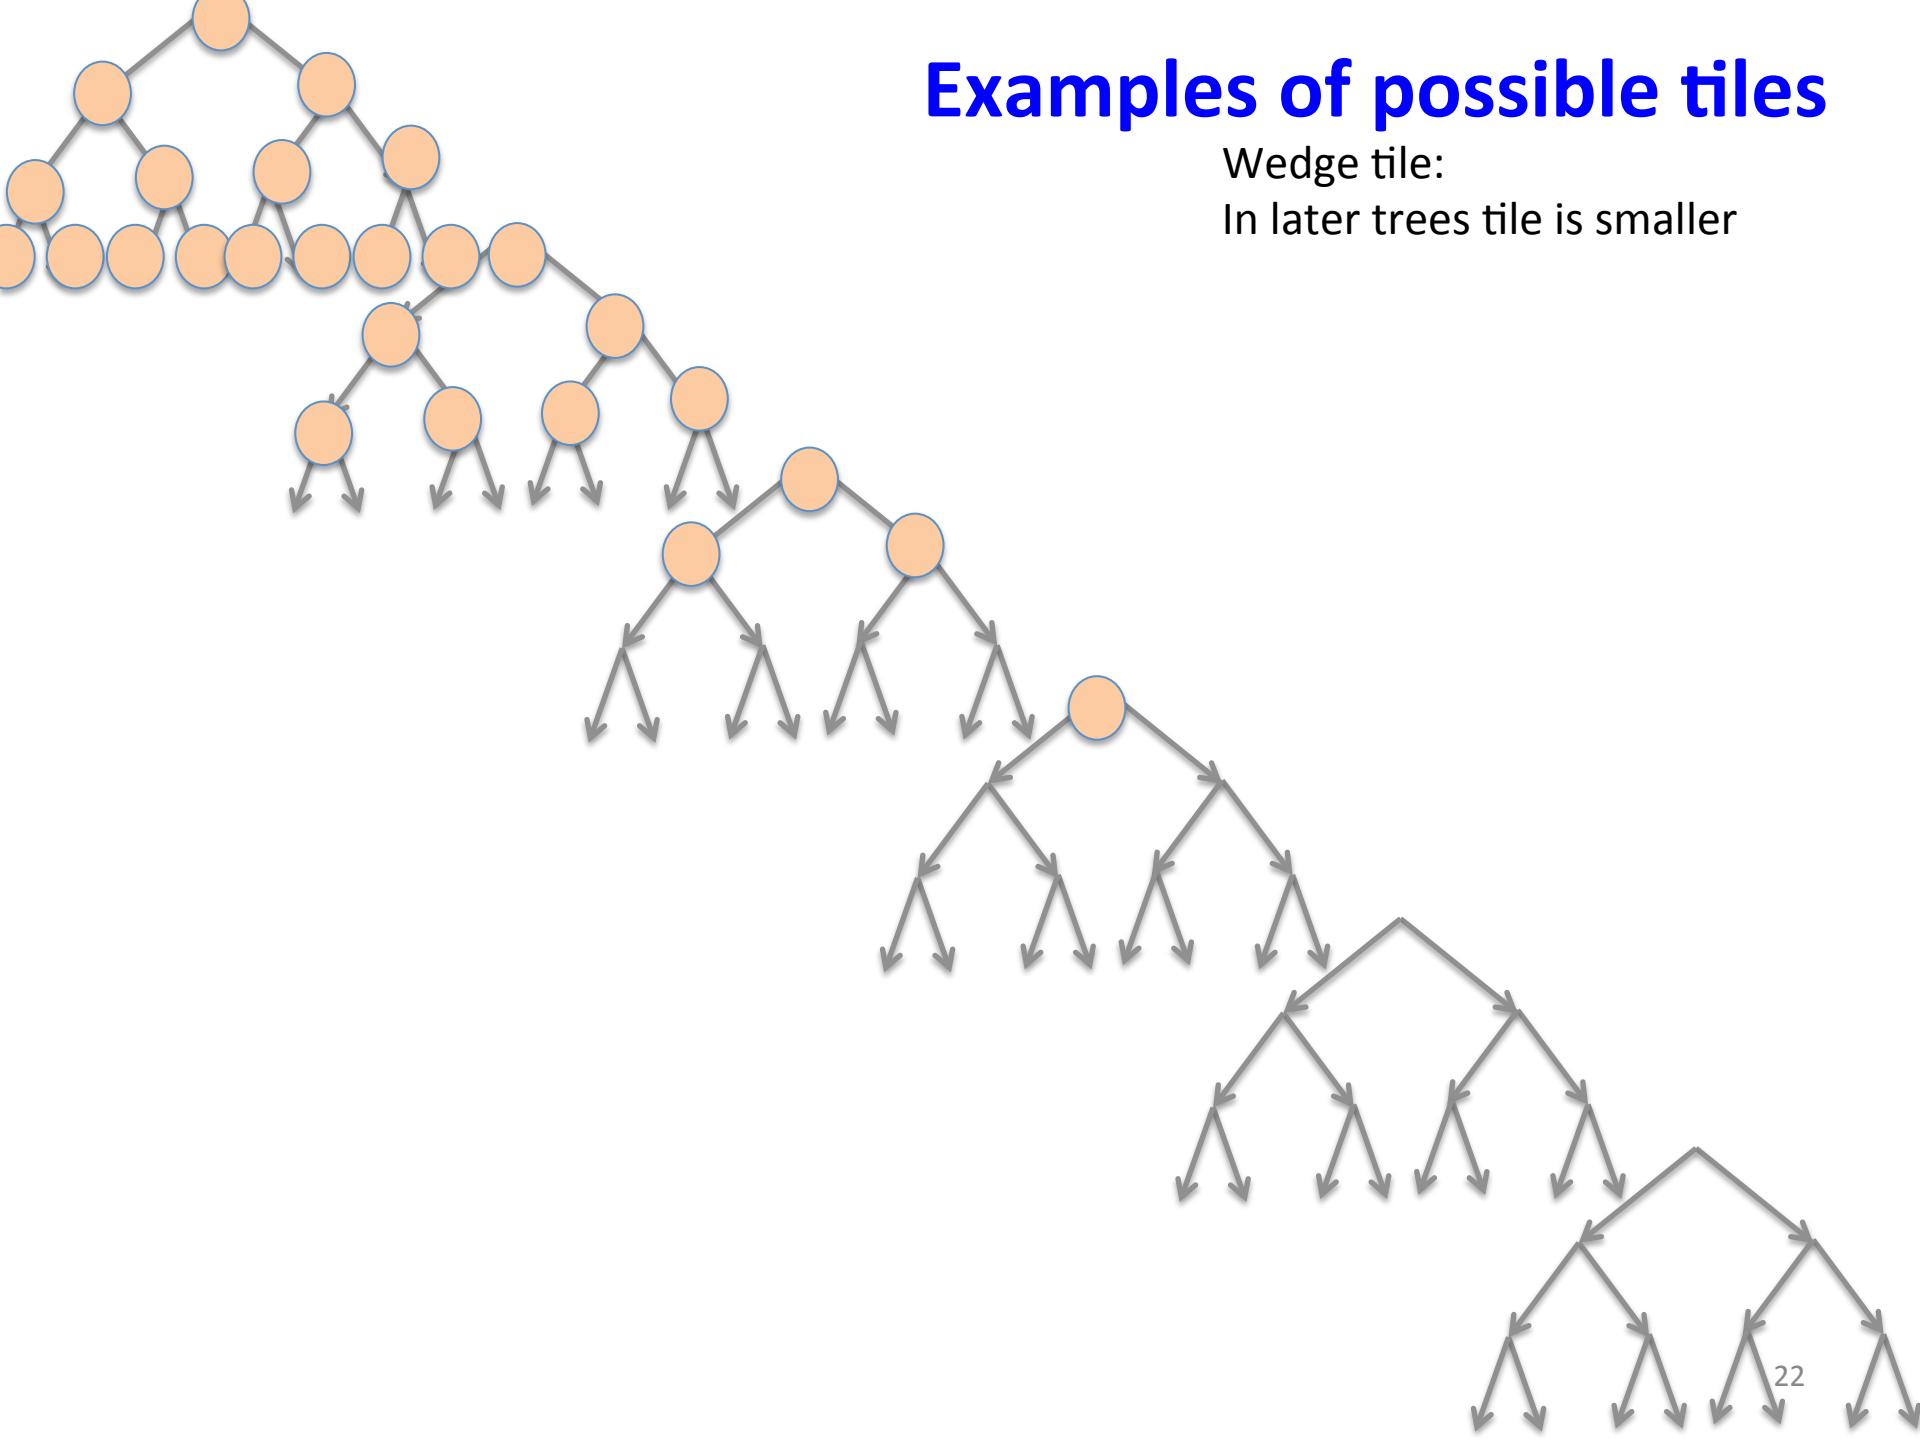
# Examples of possible tiles



# Examples of possible tiles

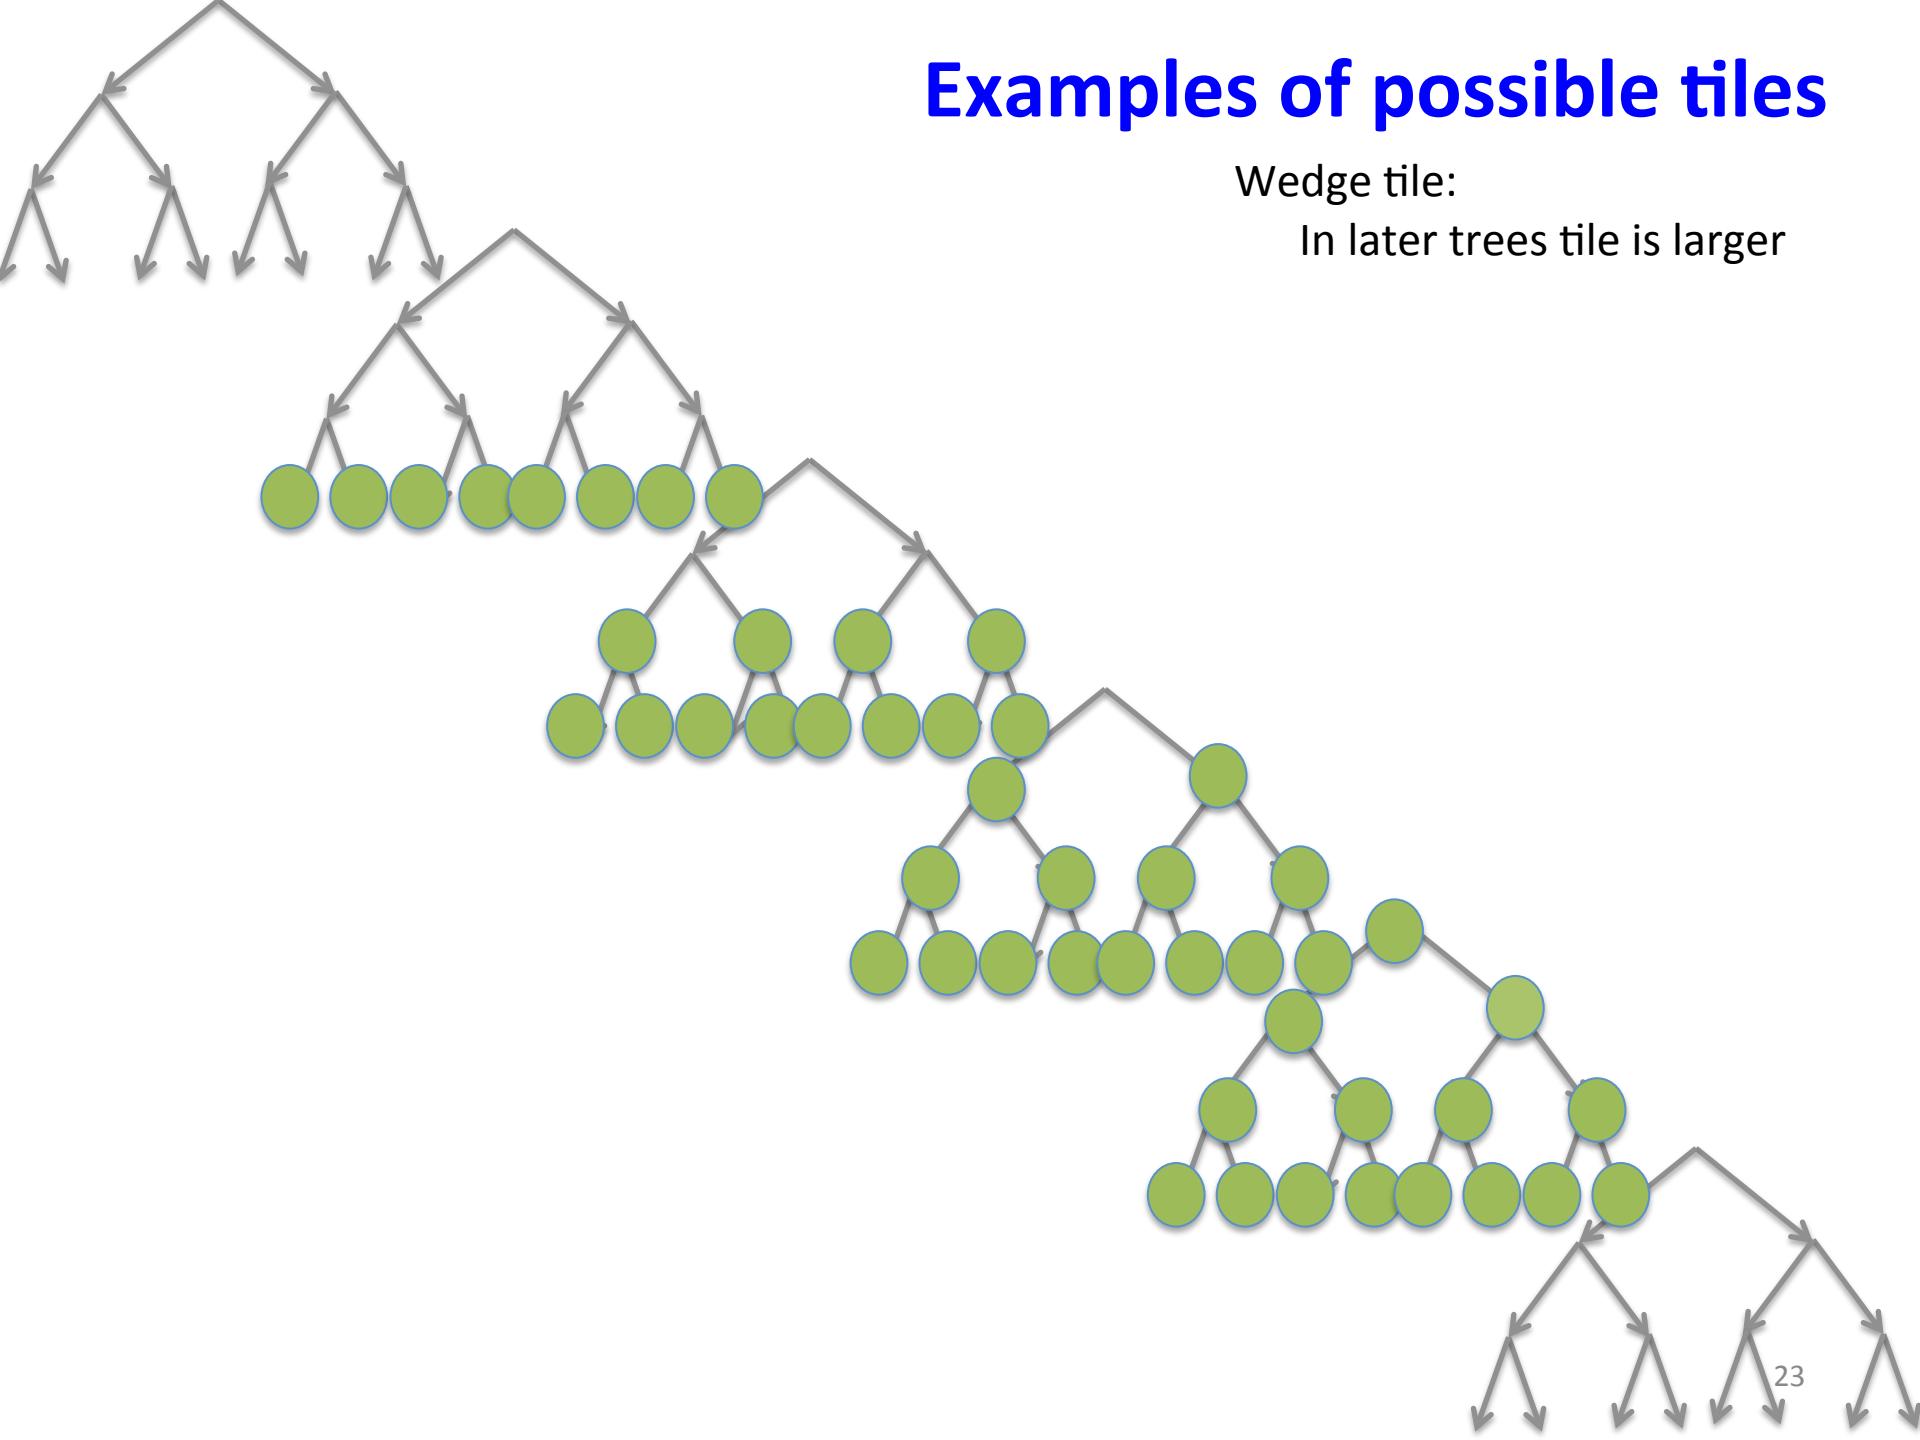
Wedge tile:

In later trees tile is smaller



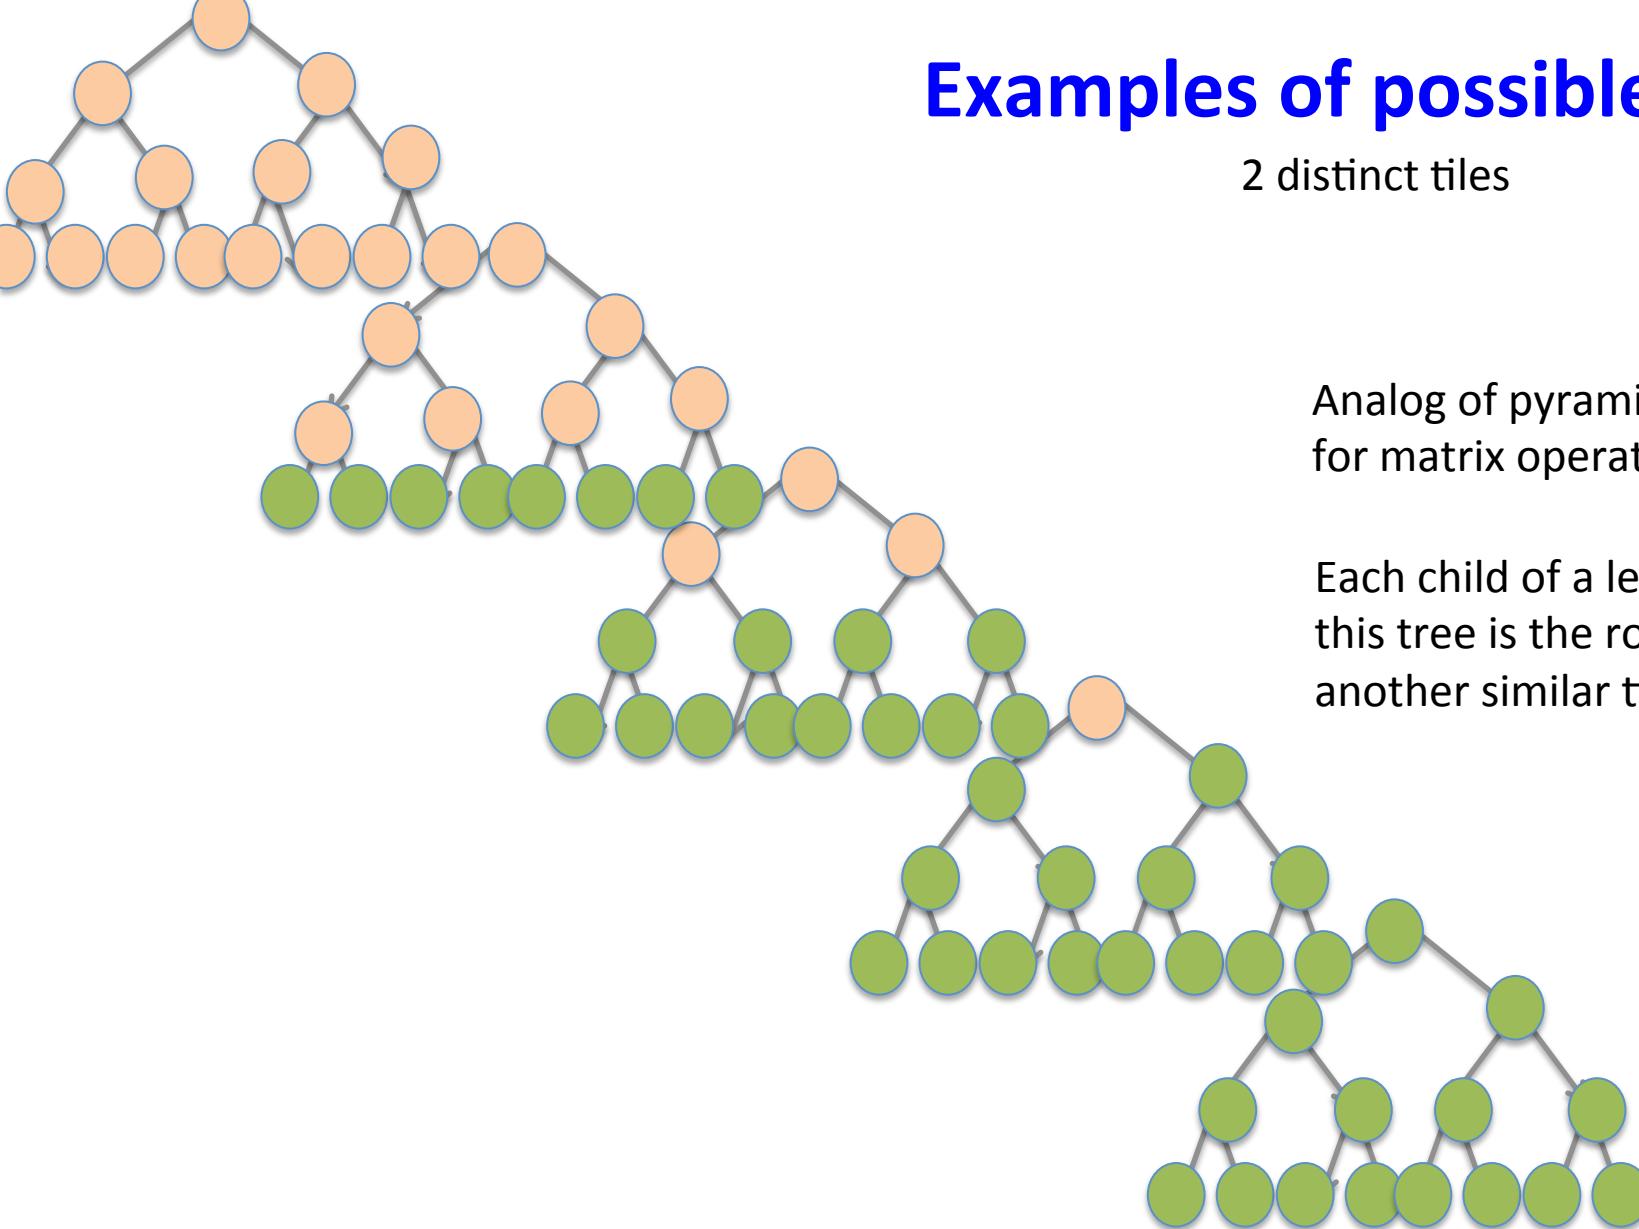
# Examples of possible tiles

Wedge tile:  
In later trees tile is larger



# Examples of possible tiles

2 distinct tiles



Analog of pyramid tiles  
for matrix operations

Each child of a leaf node in  
this tree is the root of  
another similar tiling.

# For SPX

- Develop the role of community
- Revisit array-based apps with the notion of communities in mind
  - Maybe we should look (briefly) at one other community.
  - Generalizing from 2 cases is a bit shaky.
- Develop a CnC tree-based community
  - Assess existing state
  - if necessary improve it
    - Routines for the tree-base dependence functions: parent, first-child, next-tree etc.
    - A set of tree-specific analyses, optimizations and runtimes for CnC
    - result: improve portability within the tree-base community
- Develop experience with the idea of a tree-based CnC community
  - Building applications in this class
  - Developing optimizations for this community
  - Developing analyses for this community
  - Developing runtimes for this community
- Assess the idea of role of community

# Conversations

- Now
  - ⇒ Conversation: About tuning a specific app
  - Applications expert & tuning expert
- New
  - ⇒ Two conversations
    - Community expert, applications expert & tuning expert
      - One about tuning a specific app
        - Application expert & tuning expert in context of a given community
      - One about tuning support for a specific community of apps
        - Community expert, application expert & tuning expert

# Summary

- 3 distinct roles
- A range of ways it plays out
  - One person plays all three roles
  - Two roles + a negotiation between the domain expert and the tuning expert
  - 3 people / 3 roles
  - Community + two roles (the current array case)
- This perspective isolates the work of the community & simplifies the roles of the
  - The application developer
  - The tuning expert
- It provides a capability for the domain community to define a set of dependence functions and computations that
  - Is intuitive for the application expert
  - Can be used effectively by the tuning expert to optimize domain-specific applications.

# Conclusion

- CnC doesn't know about arrays or trees but  
It does know about dependence functions
- Introduced new role: community
- A community shares a specific set of
  - Dependence functions
  - Structures
  - Analyses
  - Optimizations
  - Runtimes
- Examples:
  - Array-based community
  - Tree-based community
  - Other suggestions welcome