



National Science Foundation  
WHERE DISCOVERIES BEGIN

# *Towards Automatic Regularity Detection in Intel CnC C++*

**Louis-Noël Pouchet<sup>1</sup> and Gabriel Rodriguez<sup>2</sup>**

<sup>1</sup>Colorado State University

<sup>2</sup>University A Coruna, Spain

**October 14<sup>th</sup>, 2017**

**9<sup>th</sup> Concurrent Collections Workshop**



# 1-Slide Overview

- ◆ **Objective:** enable polyhedral optimization on (sub-)graphs which are regular/affine
  - Exploit explicit, implicit/hidden, and data-dependent regularity
- ◆ **Constraints:**
  1. Operate on C++ Intel CnC programs, but without building a C++ code analyzer
  2. Do not modify the user code: optimization is transparent to the user
  3. Generated transformed code which is always valid, whatever the input data
- ◆ **Approach:**
  1. Generate an execution trace of the program
  2. Reconstruct affine regions with specialized trace compression technique
  3. Optimize affine regions with PoCC, generate new CnC sub-graph
  4. Modified runtime: executes normal graph + affine graph (runtime skips a step in “normal” if it is already included in “affine”)

# **Motivation(s) of This Work**

**Key idea: some graphs have regularity,  
exploit it to enable static compiler optimizations**

- ◆ **Motivation (official):** enable polyhedral compilation on Intel CnC C++ graphs
- ◆ **Motivation (*in reality*):** determining when/where we can conveniently find regularity in the tag functions, without static analysis of the graph/tag functions themselves
- ◆ **Motivation (*unofficial*):** outline a system that could help detect regular sub-regions in irregular applications (e.g., MADNESS)

=> Although still preliminary, initial results show high potential ☺

# The Concept of Regularity: Purely Static

```
env::(MM:0..N,0..N,0..N);  
[A:i,k],[B:k,j],[C:i,j,k-1] -> (MM:i,j,k) -> [C:i,j,k];
```



Static analysis - model the graph as polyhedra:

```
MM : { MM[i,j,k] : 0 <= i,j,k < N };  
Reads_MM : { MM[i,j,k] -> A[i,k], B[k,j], C[i,j,k-1] };  
Writes_MM : { MM[i,j,k] -> C[i,j,k] };
```



Compile-time optimization: generate transformed polyhedral graph

```
MM_opt : { MM[ii,jj,kk] : 0 <= ii,jj,kk < N/T };  
... (tiled graph) ...
```



Compile-time code generation: produce Intel CnC C++ program from polyhedral graph

```
for(int i = 0; i < num_blocks; i++)  
    for(int j = 0; j < num_blocks; j++)  
    {  
        std::shared_ptr<Tile2d<float> > tile;  
        Triple tag = Triple(i,j,num_blocks);  
        int block_size = c.block_size;  
        c.mat_C_blocks.get(tag, tile);  
        .....  
    }
```

# The Concept of Regularity: Dynamic Discovery

```
env:: (MMsome-range) ;  
[A:tagfunc1 () ] , [B:tagfunc2 () ] , C[B:tagfunc3 () ] ->  
(MM:tagfunc4 () ) -> [C:tagfunc5 () ] ;
```



Static analysis to model the graph as polyhedra: not possible, the graph is not affine!



Runtime execution: profile the tag values generated

```
[A:0] , [B:0] , [C:0] -> (MM : 0) -> [C:1]  
[A:1] , [B:1] , [C:1] -> (MM : 1) -> [C:2]  
...  
[A:1024] , [B:1024] , [C:1024] -> (MM : 1024) -> [C:1025]  
...
```



Affine trace compression: rebuild polyhedra from trace elements

```
MM : { MM[i,j,k] : 0 <= i,j,k < N } ;  
Reads_MM : { MM[i,j,k] -> A[i,k] , B[k,j] , C[i,j,k-1] } ;  
Writes_MM : { MM[i,j,k] -> C[i,j,k] } ;
```



Compile-time optimization: generate transformed polyhedral graph

...

# **Dynamic Regularity: Pros and Cons [1/2]**

## **Pros**

### **1. Does not need any static analysis of the input program**

- Can be deeply templated Intel CnC C++ code,
- Truly, entirely independent from how the CnC program is written

### **2. Can find regular regions inside irregular programs**

- Typical example: representing a regular grid using an array of coordinates
- Can find partial regularity: a regular sub-region in the full program
- Can find “unknown” regularity: higher-dimensional regularity vs. low-dimensional irregularity

### **3. Enables full compatibility with existing polyhedral tools for CnC**

- E.g., PIPES, PoCC-DFGR, and new tools to be developed!

# **Dynamic Regularity: Pros and Cons [2/2]**

## **Cons (challenges to be solved)**

### **1. Affine trace compression is challenging**

- No unique way to represent the program, failure is very expensive
- Note: massive progresses by G. Rodriguez (CGO'16), making this work possible!

### **2. Requires to execute the original graph**

- Analysis/optimization driven by the input data set
- Highly dependent on the tag semantics implemented by the user!
- Need to ensure the transformed program remains valid for any input data!

### **3. Partial regularity may be useless**

- Finding 10 regions of one step instance each is useless, we want 1 region of 10 instances!
- No guarantee there will be any regularity when executing on new data

# Affine Trace Compression

Starting point: Rodriguez et al., “[Trace-based affine reconstruction of codes](#)”, CGO’16

- ◆ Prior work: from the trace of memory addresses accessed, rebuild the polyhedron modeling all these unique addresses
  - Super fast! (seconds for billions of entries)
  - Does not rebuild a polyhedral representation of the program
- ◆ New developments for this work:
  - Rebuild the domain (i.e., description of tag values) for steps and items
  - Connect item tags with step tags to form dataflow relation
- ◆ Key opportunities of using trace compression with CnC:
  - Data is single assignment, tags are necessarily unique
  - No need to rebuild the schedule: we can sort the tag values to improve reconstruction

# Affine Trace Compression for CnC: Status

- ◆ **Works well for the tested examples (some iCnC samples)**
  - Very fast
  - Sample apps are conveniently written with multidimensional tags
- ◆ **But potential scalability issues in later stages (poly. transformation)**
  - Rebuilt domains may contain large integer coefficients (e.g.,  $10000i+100j+k$ )
  - Need to investigate de-linearization techniques
- ◆ **And potential scalability issues for partial regularity**
  - Trace compression can always succeed, by building one polyhedron per point
  - Key difficulty: when to terminate the reconstruction in case of failure
- ◆ **Likely, need to design filtering/sorting heuristics on the input trace**
  - As CnC graph is schedule-independent, can play with sorting/filtering prior to trace compression

# Runtime Modifications

Main objective: no modification of the user code  
=> in turn, we modify the runtime ☺

- ◆ Gather graph execution trace: use iCnC tracing capabilities

```
std::ostream & cnc_format( std::ostream& os, const halo_tag & t ) {  
    os << "(" << t.t << "," << t.x << "," << t.y << "," << t.z << "," << t.f <<  
    "," << t.d << ")";  
    return os;  
}
```

- ◆ Execute transformed graph: hook into step prescription

- Main idea: generate a function **checkIsInPolyGraph(step name, tag value)** which returns true if this tag value is part of the polyhedral graph
- At start, the entire polyhedral sub-graph is prescribed
- Then the user graph/code proceeds normally
- Each time a user-code step is prescribed, if **checkIsInPolyGraph(step,tag)=true** then the step is not prescribed (it was already prescribed by the polyhedral sub-graph)

# Recommendations

---

- ◆ **Generating trace with multidimensional tags is always better**
  - Propose, natively as part of the default data structures, MULTIDIMENSIONAL INTEGER TAG CLASSES, printable
  - Right now, the user defines and implements her own tag class
  - If the classes are part of iCnC, much easier to specialize runtime code for specific tag types
- ◆ **The step/item collection names need to be printed in the trace**
  - Printer functions available, but again need to be defined by the user
- ◆ **Hooking into the prescribe function quite dirty**
  - Offer a tuner to “bypass” the prescription of a particular tag?
- ◆ **And what about OCR?**
  - These ideas apply too! ☺

# **Current Results and Status**

---

- ◆ **We only evaluated samples from the iCnC distribution**
  - Can successfully rebuild a polyhedral representation for (nearly) the full program for rtm\_stencil (halo and tiled!), sor, matrix\_inverse, heat\_equation, etc.
  - Dataset sizes are small, so “failure” of trace compression not an issue
  - Trace generation + polyhedron reconstruction is nearly automated (small manual steps)
- ◆ **We prototyped the prescribe hook for one case (manually)**
  - Polyhedron inclusion test is straightforward
  - Seems to work, but not heavily tested...
- ◆ **We did not evaluate the benefit of transformed graphs via PoCC**
  - Main issue: for good coarsening, data coarsening should be applied => user code change
  - We expect benefits shown in DFGR and PIPES work to hold
- ◆ **We still have to design a good algorithm for sub-region detection**
  - Precisely: failing “quickly enough” when a tag cannot be easily added to a polyhedron

# **Conclusion and Future Work**

---

- ♦ **Dynamic Regularity in CnC graph can be exploited**
  - Hybrid dynamic/static approach: profile once, transform, and generate always-correct code. No inspector/executor used in this work.
  - Possible only thanks to recent progresses in affine trace compression
  - Runtime modifications were minimal, approach independent from the user code
  - Preliminary results showed some of the potential of the approach, more tests needed
- ♦ **CnC + affine trace compression = good fit!**
  - CnC graphs are schedule-independent, and tag values are unique ☺
  - Still, quite some modifications/extensions needed from original CGO'16
- ♦ **Risks of this approach / limitations**
  - Totally dependent on the semantics of tags implemented by the user!
  - Totally optimistic: when executing with different data, possibly no use of opt. graph