

# Tessellating Stencils

**Liang Yuan**, Yunquan Zhang, Peng Guo, Shan Huang  
SKL of Computer Architecture, ICT, CAS

# Outline

- Introduction
- Related work
- Tessellating Stencils

# Stencil Overview

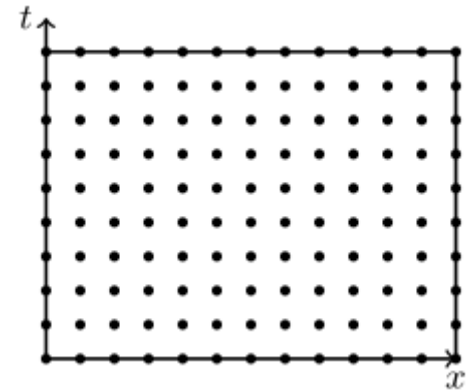
- Stencil
  - update each point in a  $d$ -dimensional space (*data space*) with a pre-defined pattern of neighbor points
  - time-iterated updates (*iteration space*)



1d3p stencil



Data space



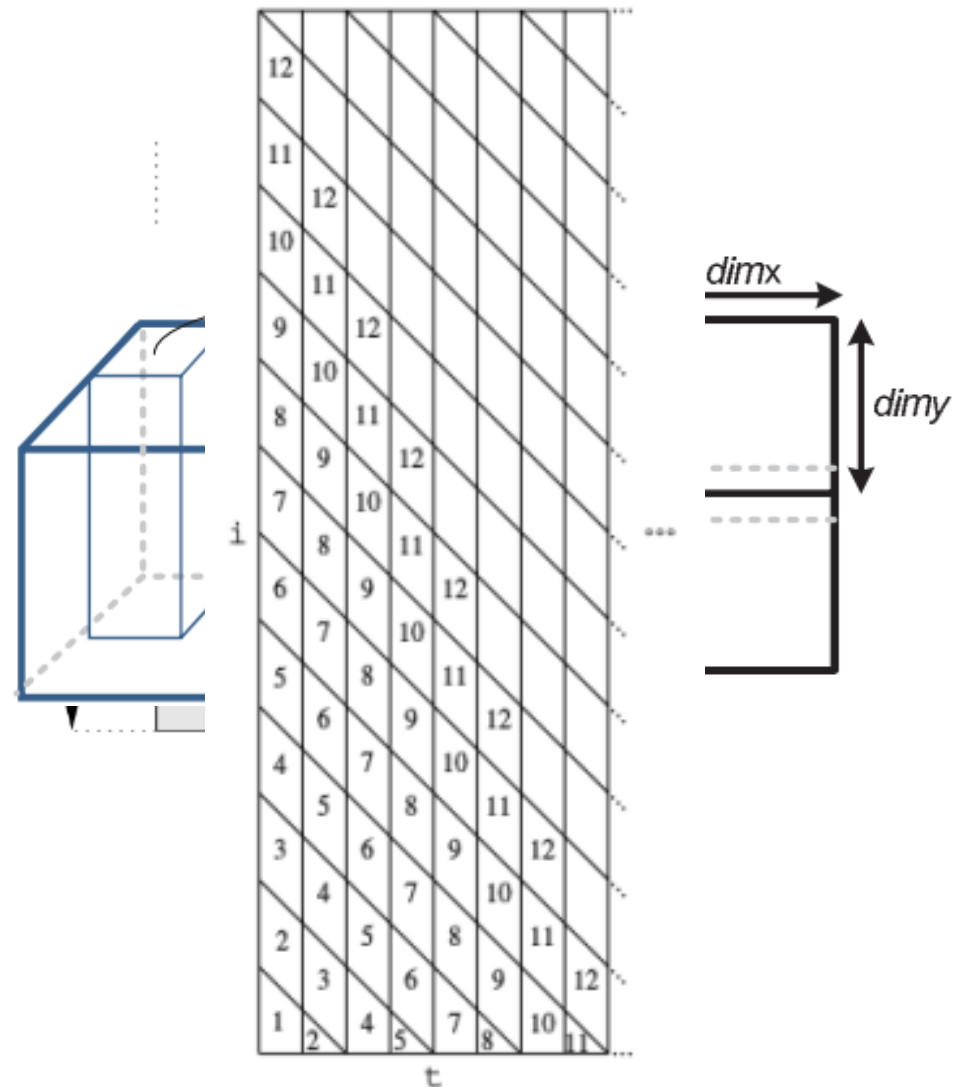
Iteration space

# Stencil Overview

- Classification
  - grid dimensions (1D, 2D, ...)
  - number of neighbors (3-point, 5-point, ...)
  - shapes (box, star, ...)
  - dependence types (Gauss-Seidel, Jacobi)
  - boundary conditions (constant, periodic, ...)

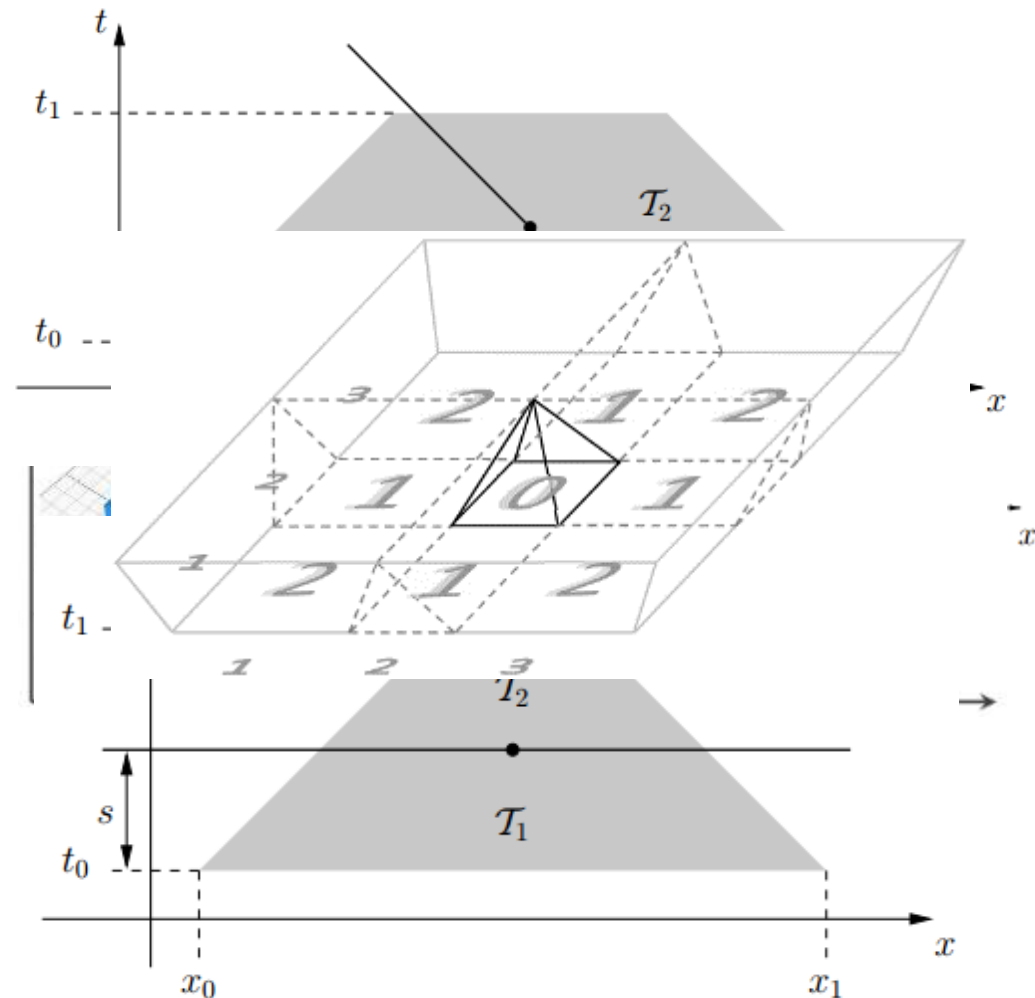
# Related Work

- Overlapped tiling
  - Hyper-rectangle [PLDI'07]
  - 3.5d blocking [SC'10]
- Time skewing
  - eliminates the redundant computations [SC'01]
  - pipelined startup and limited concurrency [SPAA'01]



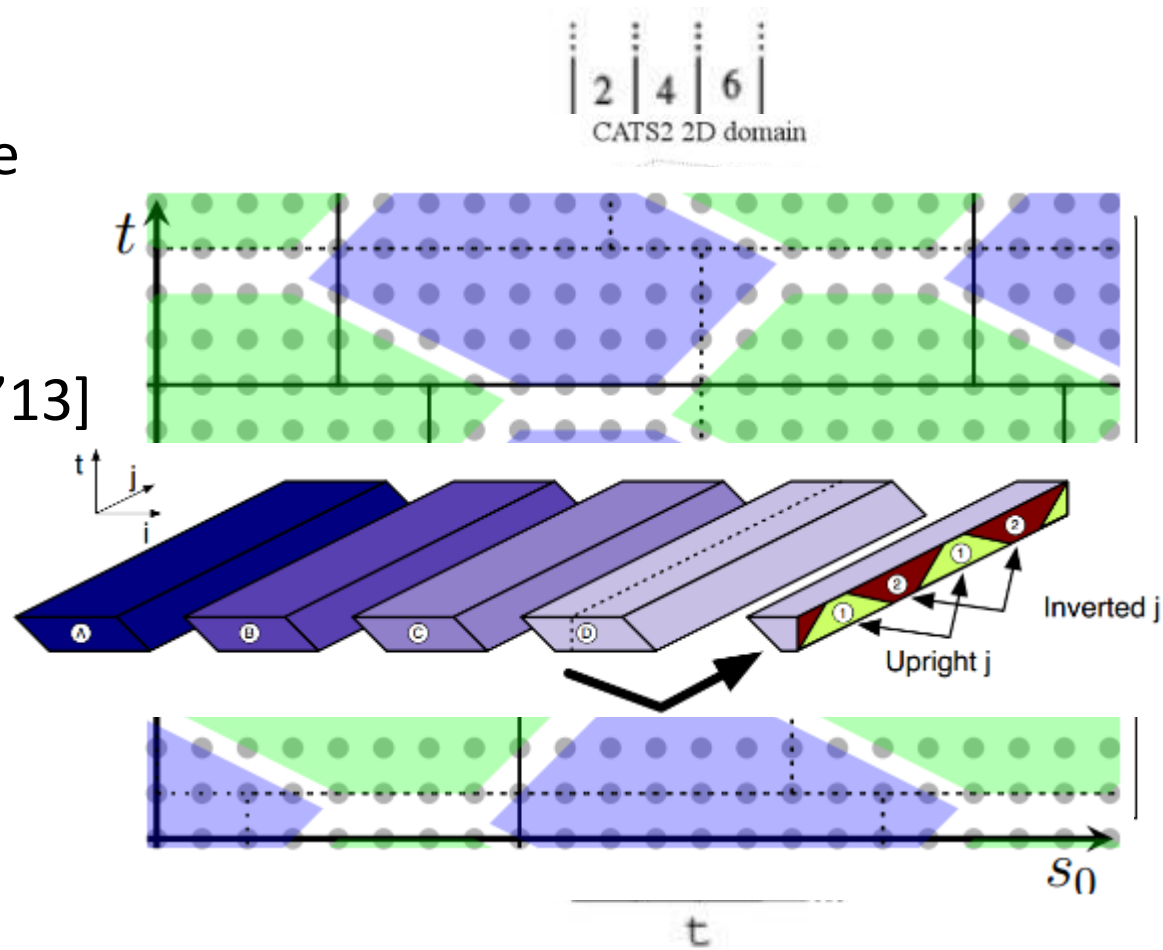
# Related Work

- Diamond tiling
  - classic block [PLDI'08]
  - higher dimensional [SC'12]
  - hexagon in 2D and truncated octahedron in 3D [PPL'14]
- Cache oblivious tiling
  - serial cache oblivious stencil algorithms [ICS'05]
  - parallel cache oblivious stencil algorithms [SPAA'06]
  - Pochoir [SPAA'11]

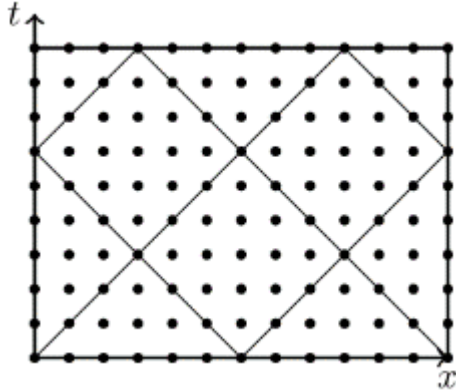


# Related Work

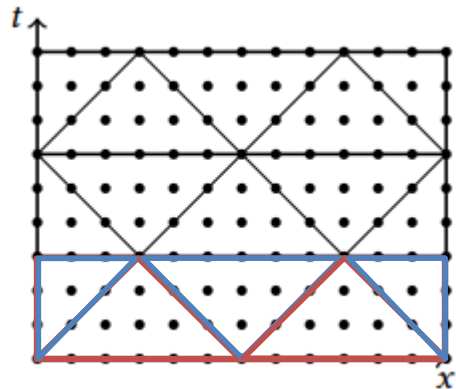
- Split tiling
  - avoids the overhead of the pipelined start-up in wavefront parallelization
  - classic split tiling [IMPACT'13]
  - nested split-tiling [ICS'13]
- Hybrid tiling
  - CATS [ICPP'11]
  - MWD [SIAMJOSC'15]
  - Grosser [CGO'14]
  - Hybrid split-tiling [ICS'13]



# Reformulating Classic 1D Diamond Tiling



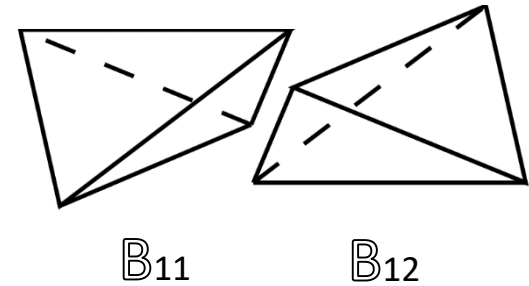
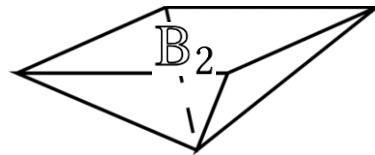
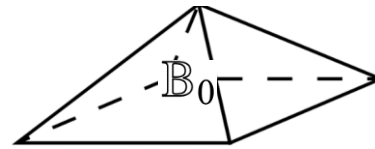
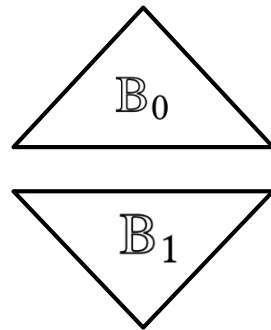
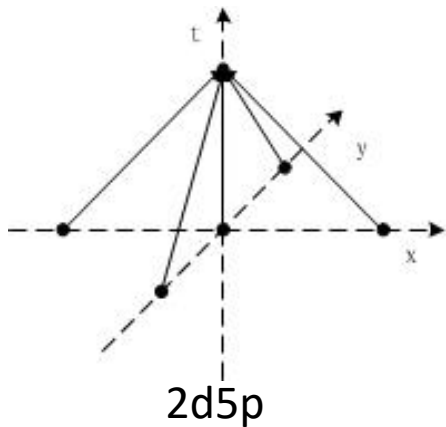
- Iteration space of 1D stencil
- Diamond tiling



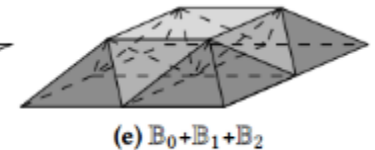
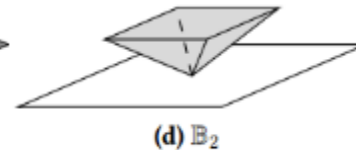
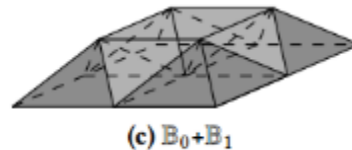
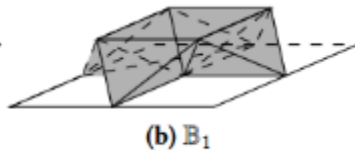
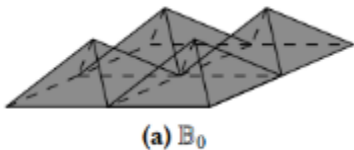
- Each diamond is further split by a horizontal line
- 2D *time tile*.
  - The region between two splitting lines
  - every grid point in a *time tile* starts in a same time dimension and is updated by identical steps.
  - interleaved computations of triangle ( $\mathbb{B}_0$ ) or inverted triangle ( $\mathbb{B}_1$ )



# Extend to 2D Stencil

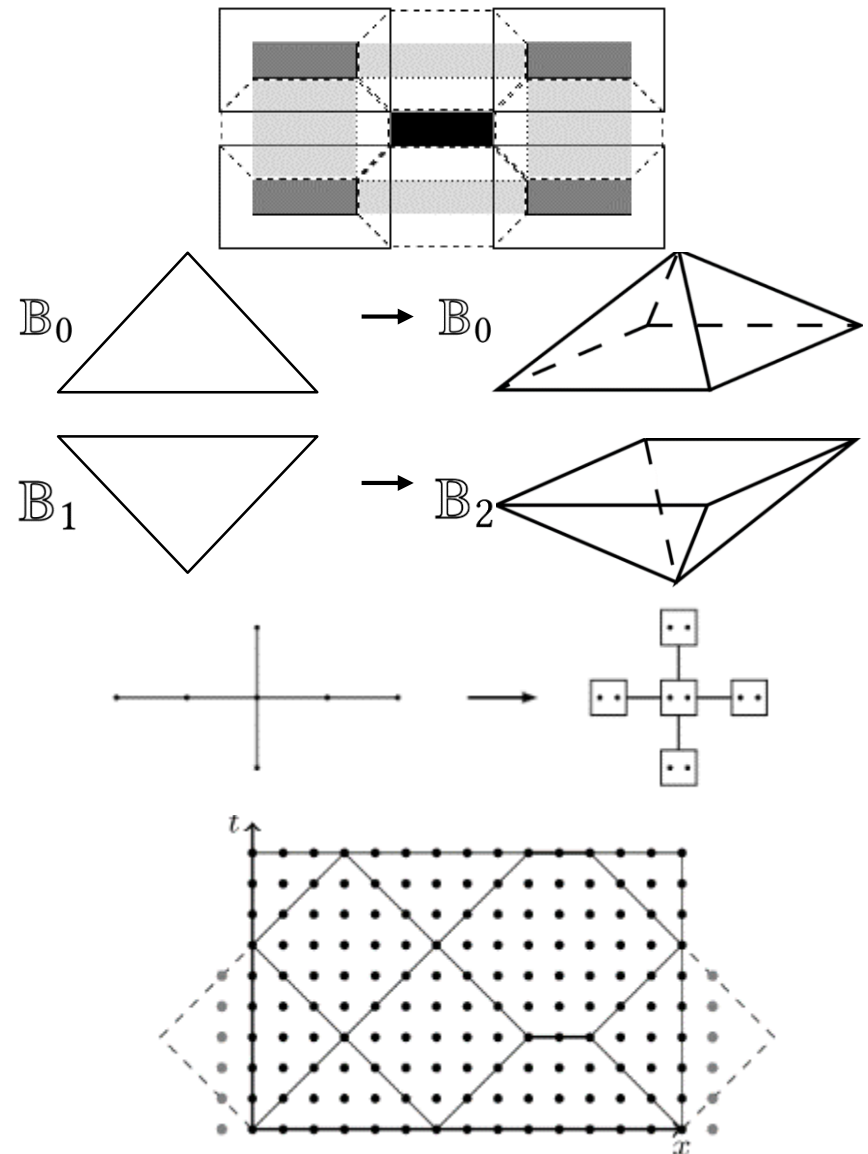


- Triangles in 1d stencil correspond to pyramids  $B_0$
- Inverted triangles correspond to inverted pyramids  $B_2$
- tetrahedrons  $B_1$



# Implementation

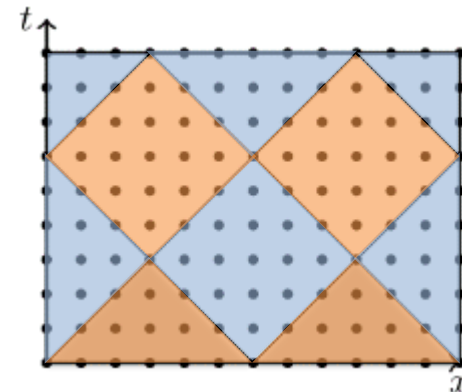
- Coarsening
  - cut in different sizes
- Merge
  - $\mathbb{B}_0$  and  $\mathbb{B}_1$  in 1D stencil.
  - $\mathbb{B}_0$  and  $\mathbb{B}_2$  in 2D stencil.
- m-order stencil
  - combine every m points to a supernode.
  - 1-order one between supernodes.
- Periodic boundary
  - stretch and transform one block



# 1D code

```
for(tt = -bt; tt < T ; tt += bt ){  
    for(n = 0; n < #B0B1[level]; n++, level = 1 - level) {  
        for(t = max(tt, 0) ; t < min( tt + 2 * bt, T); t++){  
            xmin = max( XSLOPE, xright[level] - bx + n * ix  
                        + abs(t+1, tt+bt) * XSLOPE);  
            xmax = min( N + XSLOPE, xright[level] + n * ix  
                        - abs(t+1, tt+bt) * XSLOPE);  
            for(x = xmin; x < xmax; x++){  
                update(t, x);}}}}}
```

bt: block size in time dimension  
bx: block size in data space  
bx  $\neq$  bt: support of coarsening  
#B0B1 is different in different level  
Determine the scope of 1D data space



**level=1**

**level=0**

**level=1**

**level=0**

# 2D code

```
for(tt = -bt; tt < T; tt += bt){  
    for(n = 0; n < #B0B2[level]; n++){  
        for(t = max(tt,0); t < min( tt + 2*bt, T); t++){  
            calculate xmin, xmax, ymin, ymax of B0 and B2  
            for(x = xmin; x < xmax; x++) {  
                for(y = ymin; y < ymax; y++){  
                    update(t, x, y); }  
            }  
        }  
        for(n = 0; n < #B1[0] + #B1[1]; n++){  
            for(t = tt+bt ; t < min( tt + 2*bt, T); t++) {  
                if(n<#B1[level]){  
                    calculate xmin, xmax, ymin, ymax of B11}  
                else{  
                    calculate xmin, xmax, ymin, ymax of B12}  
                for(x = xmin; x < xmax; x++) {  
                    for(y = ymin; y < ymax; y++){  
                        update(t, x, y); }  
                    }  
                }  
            }  
        }  
    }  
    level = 1 - level;}
```

#B0B2 represents #B0 in level 0 or #B2 in level 1

two kinds of B1: B11 and B12

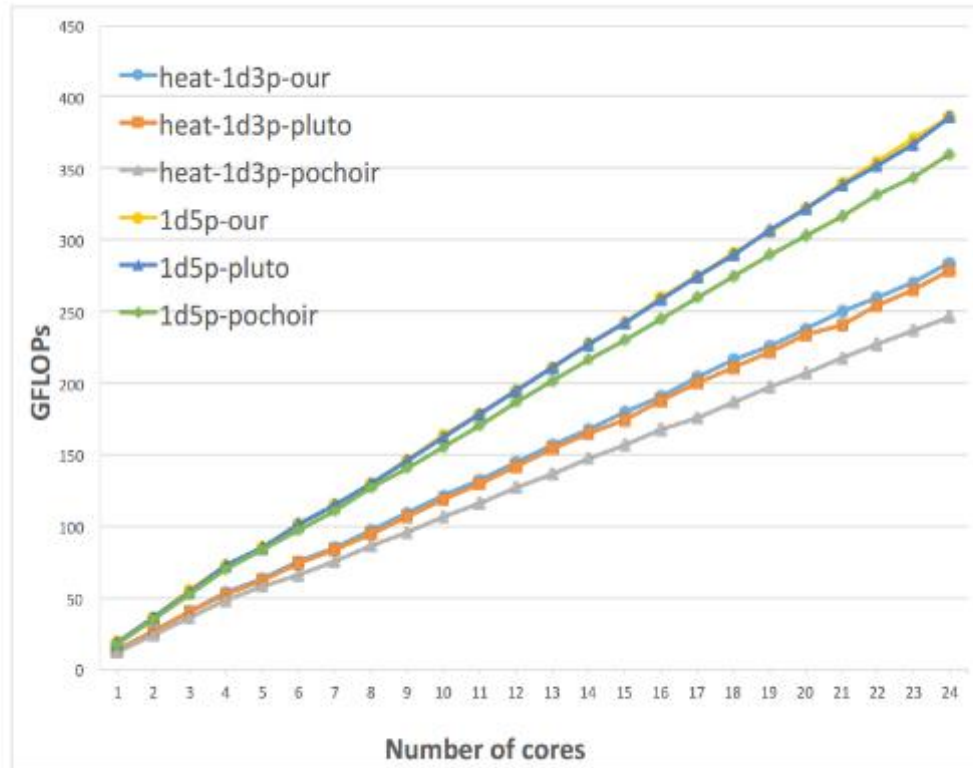
for 3D stencil we leave the unit-stride dimension uncut.

# Evaluation

- two Intel Xeon E5-2670 processors with 2.70 GHz clock speed
- 32KB private L1 cache
- 256KB private L2 cache
- a unified 30MB L3 cache shared by 12 cores
- ICC compiler version 16.0.1, flag '-O3 -openmp'.

Benchmark	Problem Size	our blocking	Pluto blocking
Heat-1D	$12000000 \times 4000$	$2000 \times 1000$	$2000 \times 2000$
1d5p	$12000000 \times 4000$	$2000 \times 500$	$2000 \times 2000$
Heat-2D	$6000^2 \times 2000$	$128 \times 256 \times 64$	$64 \times 64 \times 64$
2d9p	$6000^2 \times 2000$	$128 \times 256 \times 64$	$64 \times 64 \times 64$
Game of life	$6000^2 \times 2000$	$128 \times 256 \times 64$	$128 \times 128 \times 128$
Heat-3D	$256^3 \times 1000$	$24 \times 24 \times 12$	$12 \times 12 \times 12$
3d27p	$256^3 \times 1000$	$24 \times 24 \times 12$	$12 \times 12 \times 12$

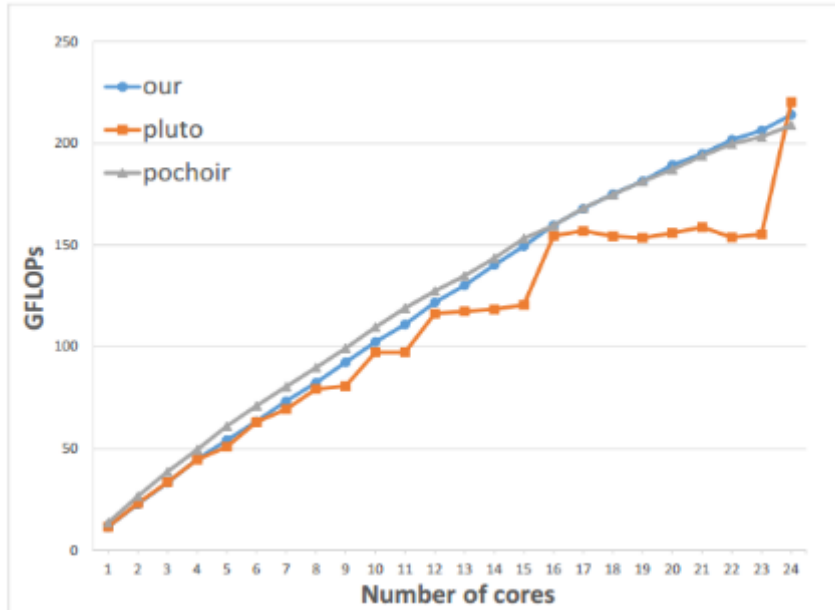
# Experimental Results



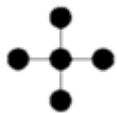
same to Pluto  
better than Pochoir

- Our scheme and PluTo produce the same diamond tiling codes
- Pochoir generates trapezoidal blocks of different sizes.

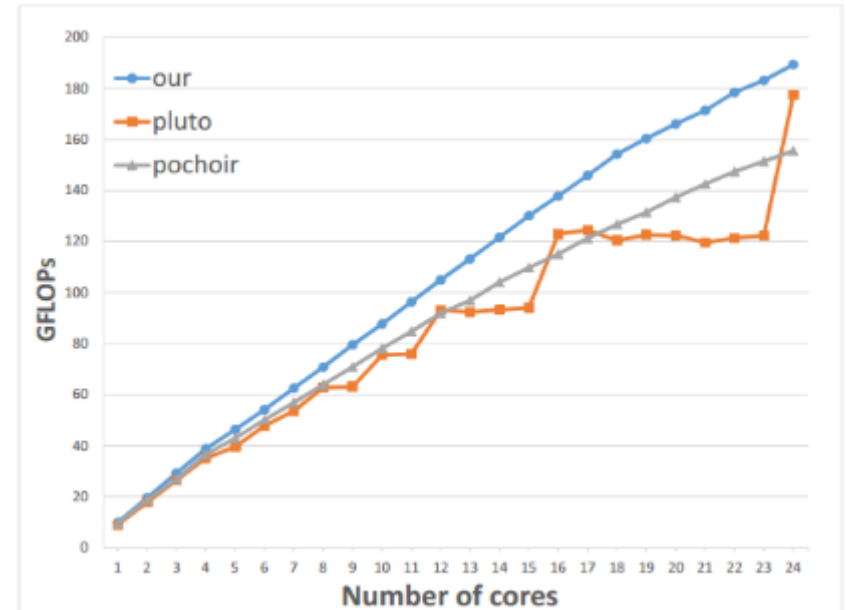
# Experimental Results



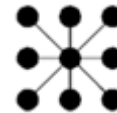
2d5p



Pluto outperforms less than 5%  
with 24 cores



2d9p

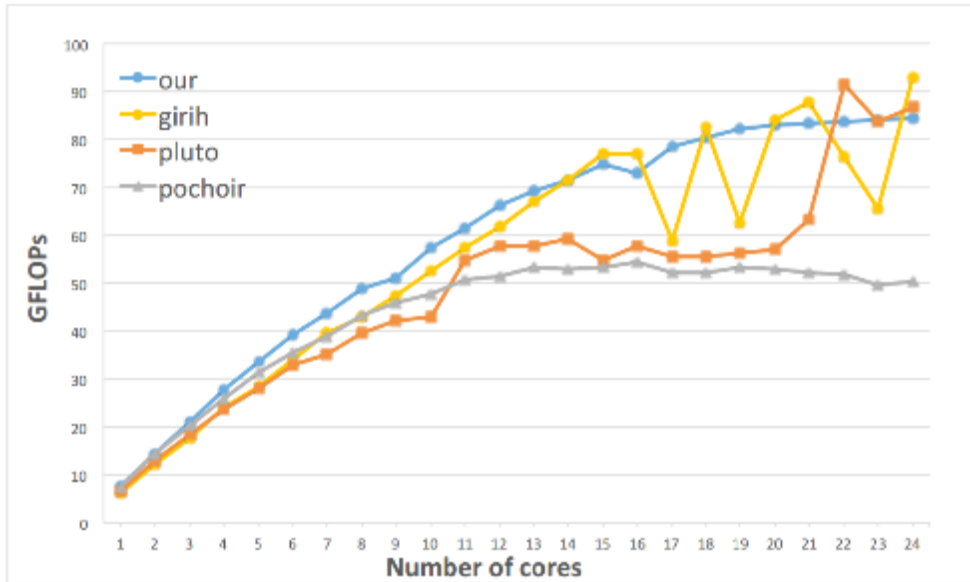


Our code outperforms  
Pluto and Pochoir  
by 14% and 20% on average.

- more suitable for box stencil.

# Experimental Results

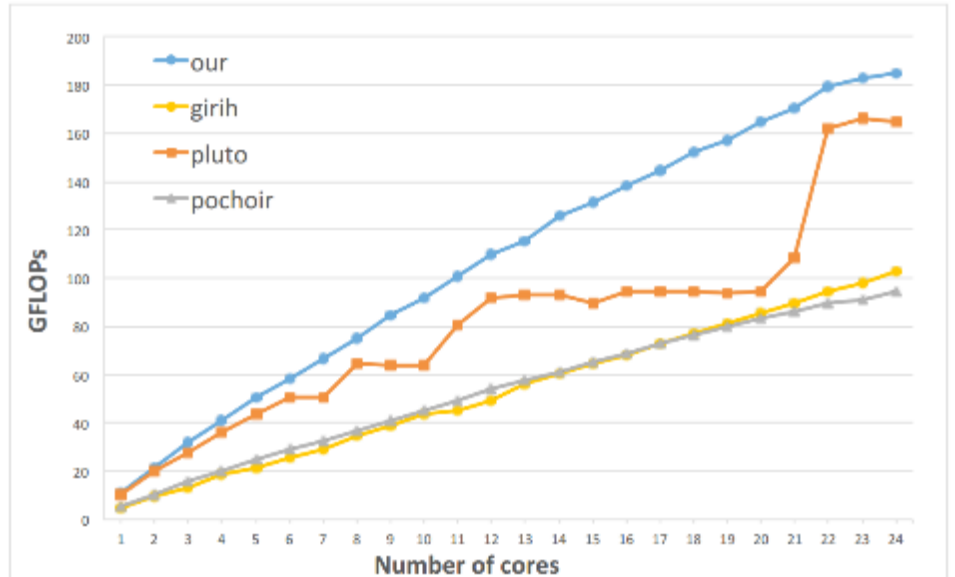
- 3D stencil leaves unit-stride uncut



3d7p

Our code and Pochoir exhibit better scalability than Pluto.

Our code outperforms Pluto and Pochoir by a maximum of **33%** and **68%**, and by **22%** and **31%** on average.



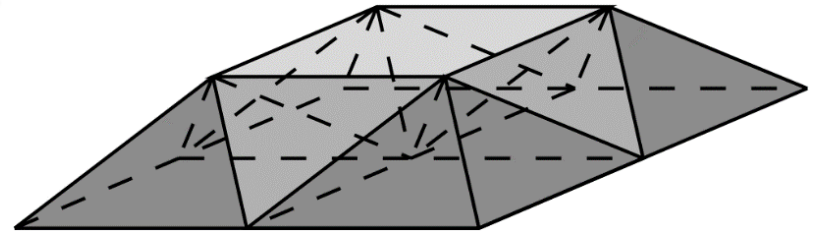
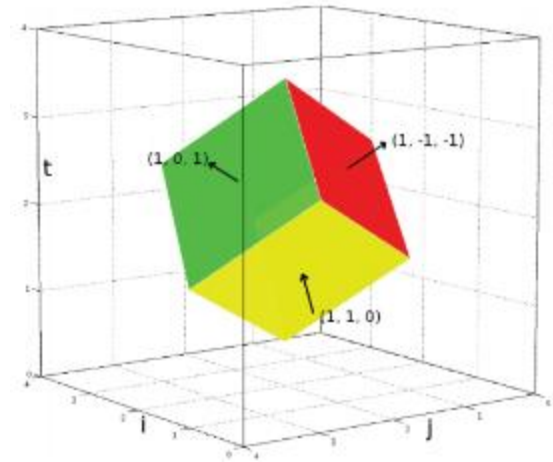
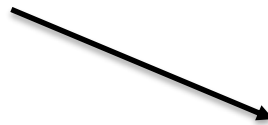
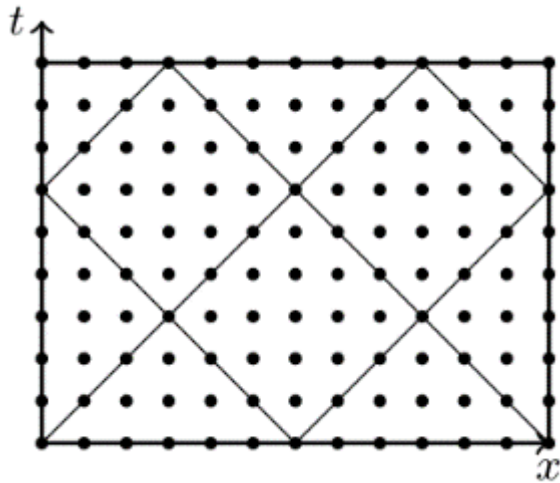
3d27p

Our code outperforms Pluto and Pochoir by a maximum of **74%** and **100%**, and by **30%** and **99%** on average.



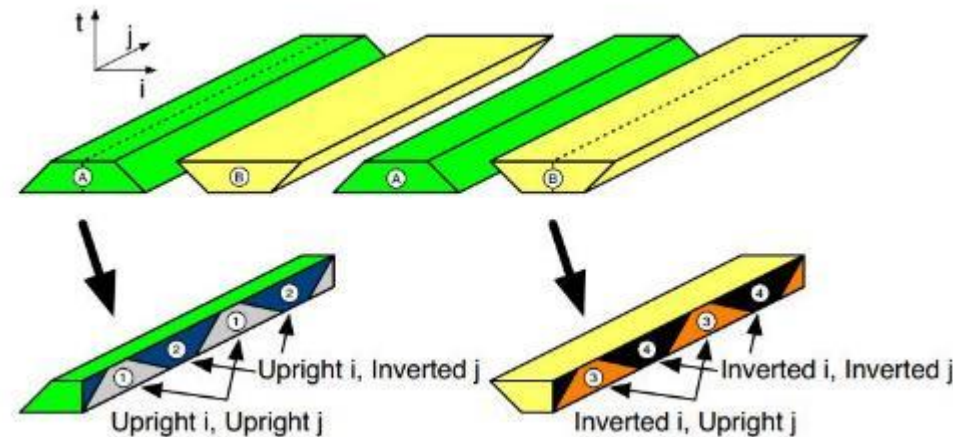
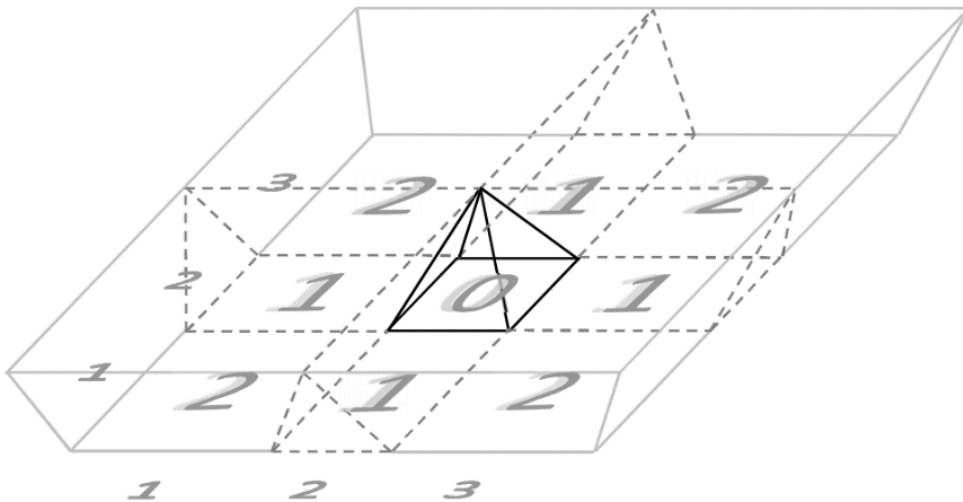
# new?

- extension of 1D diamond



# Not that new!

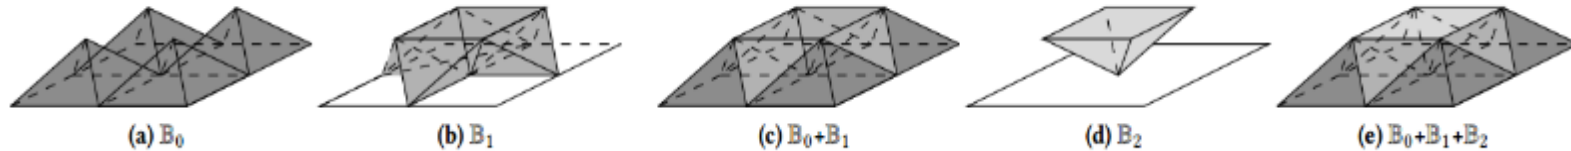
- Existing techniques produce similar blocks
  - Cache oblivious [SPAA'11]
  - Nested split-tiling [ICS'13]
  - divide multiple dimensions simultaneously



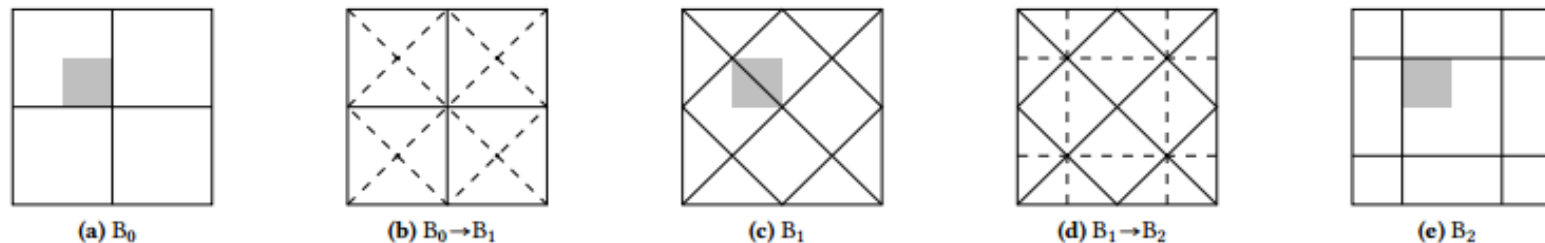
# Advantage!

- Diamond tiling
  - fixed tile size at compile time
  - small size at the apex of a diamond
  - hard to choose the proper tile sizes to ensure the concurrent start
- Cache oblivious tiling
  - overhead of recursion
  - artificial dependencies
- Split tiling
  - synchronization overhead  $2^d$

# Formulating the two-level tessellation



Project  $B_i$  in iteration space on data space, denoted as  $B_i$

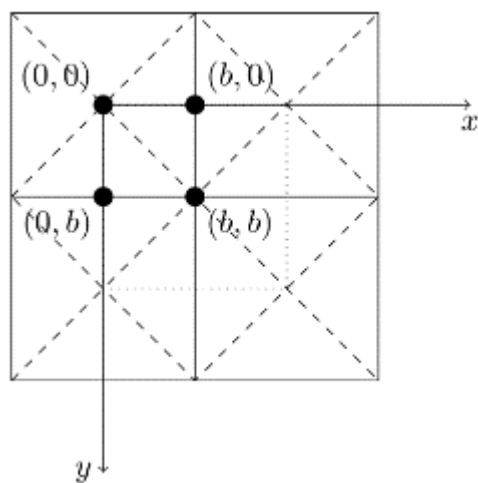


Determine the number of updating of each point in  $B_i$ .

	$B_i^j$	step1	step2	step3	$B_i$
$i = 0$	<pre> 0 0 0 0 0 0 0 0 1 1 1 1 1 0 0 1 2 2 2 2 0 0 1 2 3 3 2 1 0 0 1 2 3 3 2 1 0 0 1 2 3 3 2 1 0 0 1 1 1 1 1 0 0 0 0 0 0 0 0                     </pre>	<pre> 1                     </pre>	<pre> 1 1 1 1 1 1 1 1                     </pre>	1	
$i = 1$	<pre> 0 0 1 0 0 0 2 1 0 0 1 2 3 2 1 0 0 1 2 3 2 1 0 0 1 0 0                     </pre>	<pre> 1 1 1 1 1                     </pre>	<pre> 1 1 1 1 1 1 1 1 1 1 1 1                     </pre>	1 1 1 1 1	
$i = 1$	<pre> 0 0 1 0 0 1 2 1 0 0 1 2 3 2 1 0 0 1 2 3 2 1 0 0 1 0 0                     </pre>	1 1 1 1 1	<pre> 1 1 1 1 1 1 1 1 1 1 1 1                     </pre>	<pre> 1 1 1 1 1                     </pre>	
$i = 2$	<pre> 0 0 0 0 0 0 0 0 1 1 1 1 1 0 0 1 2 2 2 2 1 0 0 1 2 3 3 2 1 0 0 1 2 3 3 2 1 0 0 1 2 3 3 2 1 0 0 1 1 1 1 1 0 0 0 0 0 0 0 0                     </pre>	1	<pre> 1 1 1 1 1 1 1 1 1 1 1 1                     </pre>	<pre> 1                     </pre>	

# Formulating the two-level tessellation

- Consider a point  $a$  ( $a_0, \dots, a_{d-1}$ ) in  $B_0(0, \dots, 0)$
- updated time in  $B_i$ , denoted as  $T_i(a_0, \dots, a_{d-1})$ .



$$T_i(a_0, \dots, a_{d-1}) = \min(b, a_i, \dots, a_{d-1}) - \max(0, a_0, \dots, a_{i-1})$$

$$\sum_{i=0}^d T_i(a_0, \dots, a_{d-1}) = (b - a_0) + (a_0 - a_1) + \dots + (a_{d-2} - a_{d-1}) + a_{d-1} = b$$

$$a_i \geq a_{i+1}$$

# Summary

- two-level tessellation scheme
  - with highly concurrent execution
  - executes without redundant computation
  - achieves maximize parallelism
  - without false dependencies.
  - Calculate blocks without relying on the compiler
  - Extend to n-dimensional stencil
- mathematical structure of tessellation
  - Associate coordinates of elements with number of updating steps for each block.